

GTOS Detailed Volume - Volume 09 Ai Operating System

GTOS Enterprise Engineering Handbook - Final Edition 1.0

Source Chapters

- AI Operating System Introduction
- Platform AI Runtime SDK
- AI Gateway and Model Access
- AI Orchestrator Ownership
- Agent Run State Machine
- Supervisor and Delegation
- Context Builder
- Planner and Plan Versioning
- Critic and Verifier
- Policy and Approval Engine
- Typed Tool Executor
- Outcome Observer
- Recovery and Replanner
- MCP Tool Registry
- MCP Tool Gateway
- Workflow Engine Integration
- Event Bus Integration
- Provider Adapter Architecture
- Model Registry and Routing
- Prompt and Template Governance
- Structured Output and Schema Validation
- AI Memory
- RAG and Retrieval
- Knowledge Cutoff and Provenance
- Digital Twin and World Model
- Evaluation and Promotion

- Online Monitoring and Drift
- Kill Switch and Incident Containment
- AI Security and Prompt Injection
- Browser Agent Sandboxing
- Code and Data Tool Sandboxing
- AI Cost and Budget Governance
- Human Handoff
- AI Compliance Service
- AI Vision Service
- AI Voice Service
- AI NLP Service
- AI Automation Service
- AI Edge and Federated Services
- AI Marketplace
- AI Observability and AIOps
- Predictive Analytics and Forecast ML
- Fraud Engine
- Search and Ranking
- Domain AI Role Catalogue
- AI Canonical Ownership Matrix

AI Operating System Introduction

Metadata	Value
Volume	9
Chapter ID	V09-00
Subject	AIOperatingSystemIntroduction
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Operating System Introduction.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIOperatingSystemIntroduction	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOperatingSystemIntroductionVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOperatingSystemIntroductionDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOperatingSystemIntroductionEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOperatingSystemIntroductionException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIOperatingSystemIntroduction {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAI0peratingSystemIntroduction
- ValidateAI0peratingSystemIntroduction
- ApproveAI0peratingSystemIntroduction
- ActivateAI0peratingSystemIntroduction
- SuspendAI0peratingSystemIntroduction
- CorrectAI0peratingSystemIntroduction
- RetireAI0peratingSystemIntroduction

Reference Events:

- AI0peratingSystemIntroductionCreated
- AI0peratingSystemIntroductionValidated
- AI0peratingSystemIntroductionApproved
- AI0peratingSystemIntroductionActivated
- AI0peratingSystemIntroductionSuspended
- AI0peratingSystemIntroductionCorrected
- AI0peratingSystemIntroductionRetired

Reference queries:

- GetAI0peratingSystemIntroduction
- ListAI0peratingSystemIntroductionRecords
- GetAI0peratingSystemIntroductionTimeline
- GetAI0peratingSystemIntroductionEvidence
- GetAI0peratingSystemIntroductionDecisionPackage
- GetAI0peratingSystemIntroductionExceptions
- GetAI0peratingSystemIntroductionAuditTrail

Reference API surface:

```

POST /api/v1/aioperatingsystemintroduction/commands
GET  /api/v1/aioperatingsystemintroduction/{id}
GET  /api/v1/aioperatingsystemintroduction/{id}/timeline
GET  /api/v1/aioperatingsystemintroduction/{id}/evidence
GET  /api/v1/aioperatingsystemintroduction/{id}/decisions
POST /api/v1/aioperatingsystemintroduction/{id}/exceptions
POST /api/v1/aioperatingsystemintroduction/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aioperatingsystemintroduction_aggregate	authoritative subject
aioperatingsystemintroduction_version	immutable submitted, approved, issued, or signed versions
aioperatingsystemintroduction_relationship	governed relationships
aioperatingsystemintroduction_decision	binding Decisions and rationale
aioperatingsystemintroduction_evidence_link	provenance and evidence relationships
aioperatingsystemintroduction_event_outbox	transactional Event publication
aioperatingsystemintroduction_idempotency	duplicate-write protection
aioperatingsystemintroduction_exception	exception, claim, dispute, or hold
aioperatingsystemintroduction_correction	correction and supersession lineage
aioperatingsystemintroduction_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Platform AI Runtime SDK

Metadata	Value
Volume	9
Chapter ID	V09-01
Subject	PlatformAIRuntimeSDK
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Platform AI Runtime SDK.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
PlatformAIRuntimeSDK	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlatformAIRuntimeSDKVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlatformAIRuntimeSDKDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlatformAIRuntimeSDKEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlatformAIRuntimeSDKException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
PlatformAIRuntimeSDK {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreatePlatformAIRuntimeSDK
- ValidatePlatformAIRuntimeSDK
- ApprovePlatformAIRuntimeSDK
- ActivatePlatformAIRuntimeSDK
- SuspendPlatformAIRuntimeSDK
- CorrectPlatformAIRuntimeSDK
- RetirePlatformAIRuntimeSDK

Reference Events:

- PlatformAIRuntimeSDKCreated
- PlatformAIRuntimeSDKValidated
- PlatformAIRuntimeSDKApproved
- PlatformAIRuntimeSDKActivated
- PlatformAIRuntimeSDKSuspended
- PlatformAIRuntimeSDKCorrected
- PlatformAIRuntimeSDKRetired

Reference queries:

- GetPlatformAIRuntimeSDK
- ListPlatformAIRuntimeSDKRecords
- GetPlatformAIRuntimeSDKTimeline
- GetPlatformAIRuntimeSDKEvidence
- GetPlatformAIRuntimeSDKDecisionPackage
- GetPlatformAIRuntimeSDKExceptions
- GetPlatformAIRuntimeSDKAuditTrail

Reference API surface:

```

POST /api/v1/platformairtimesdk/commands
GET  /api/v1/platformairtimesdk/{id}
GET  /api/v1/platformairtimesdk/{id}/timeline
GET  /api/v1/platformairtimesdk/{id}/evidence
GET  /api/v1/platformairtimesdk/{id}/decisions
POST /api/v1/platformairtimesdk/{id}/exceptions
POST /api/v1/platformairtimesdk/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
platformairuntimesdk_aggregate	authoritative subject
platformairuntimesdk_version	immutable submitted, approved, issued, or signed versions
platformairuntimesdk_relationship	governed relationships
platformairuntimesdk_decision	binding Decisions and rationale
platformairuntimesdk_evidence_link	provenance and evidence relationships
platformairuntimesdk_event_outbox	transactional Event publication
platformairuntimesdk_idempotency	duplicate-write protection
platformairuntimesdk_exception	exception, claim, dispute, or hold
platformairuntimesdk_correction	correction and supersession lineage
platformairuntimesdk_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

AI Gateway and Model Access

Metadata	Value
Volume	9
Chapter ID	V09-02
Subject	AIGatewayandModelAccess
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Gateway and Model Access.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIGatewayandModelAccess	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIGatewayandModelAccessVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIGatewayandModelAccessDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIGatewayandModelAccessEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIGatewayandModelAccessException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

AIGatewayandModelAccess {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIGatewayandModelAccess
- ValidateAIGatewayandModelAccess
- ApproveAIGatewayandModelAccess
- ActivateAIGatewayandModelAccess
- SuspendAIGatewayandModelAccess
- CorrectAIGatewayandModelAccess
- RetireAIGatewayandModelAccess

Reference Events:

- AIGatewayandModelAccessCreated
- AIGatewayandModelAccessValidated
- AIGatewayandModelAccessApproved
- AIGatewayandModelAccessActivated
- AIGatewayandModelAccessSuspended
- AIGatewayandModelAccessCorrected
- AIGatewayandModelAccessRetired

Reference queries:

- GetAIGatewayandModelAccess
- ListAIGatewayandModelAccessRecords
- GetAIGatewayandModelAccessTimeline
- GetAIGatewayandModelAccessEvidence
- GetAIGatewayandModelAccessDecisionPackage
- GetAIGatewayandModelAccessExceptions
- GetAIGatewayandModelAccessAuditTrail

Reference API surface:

```

POST /api/v1/aigatewayandmodelaccess/commands
GET  /api/v1/aigatewayandmodelaccess/{id}
GET  /api/v1/aigatewayandmodelaccess/{id}/timeline
GET  /api/v1/aigatewayandmodelaccess/{id}/evidence
GET  /api/v1/aigatewayandmodelaccess/{id}/decisions
POST /api/v1/aigatewayandmodelaccess/{id}/exceptions
POST /api/v1/aigatewayandmodelaccess/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aigatewayandmodelaccess_aggregate	authoritative subject
aigatewayandmodelaccess_version	immutable submitted, approved, issued, or signed versions
aigatewayandmodelaccess_relationship	governed relationships
aigatewayandmodelaccess_decision	binding Decisions and rationale
aigatewayandmodelaccess_evidence_link	provenance and evidence relationships
aigatewayandmodelaccess_event_outbox	transactional Event publication
aigatewayandmodelaccess_idempotency	duplicate-write protection
aigatewayandmodelaccess_exception	exception, claim, dispute, or hold
aigatewayandmodelaccess_correction	correction and supersession lineage
aigatewayandmodelaccess_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Orchestrator Ownership

Metadata	Value
Volume	9
Chapter ID	V09-03
Subject	AIOrchestratorOwnership
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Orchestrator Ownership**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIOrchestratorOwnership	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOrchestratorOwnershipVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOrchestratorOwnershipDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOrchestratorOwnershipEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOrchestratorOwnershipException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIOrchestratorOwnership {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIOrchestratorOwnership
- ValidateAIOrchestratorOwnership
- ApproveAIOrchestratorOwnership
- ActivateAIOrchestratorOwnership
- SuspendAIOrchestratorOwnership
- CorrectAIOrchestratorOwnership
- RetireAIOrchestratorOwnership

Reference Events:

- AIOrchestratorOwnershipCreated
- AIOrchestratorOwnershipValidated
- AIOrchestratorOwnershipApproved
- AIOrchestratorOwnershipActivated
- AIOrchestratorOwnershipSuspended
- AIOrchestratorOwnershipCorrected
- AIOrchestratorOwnershipRetired

Reference queries:

- GetAIOrchestratorOwnership
- ListAIOrchestratorOwnershipRecords
- GetAIOrchestratorOwnershipTimeline
- GetAIOrchestratorOwnershipEvidence
- GetAIOrchestratorOwnershipDecisionPackage
- GetAIOrchestratorOwnershipExceptions
- GetAIOrchestratorOwnershipAuditTrail

Reference API surface:

```

POST /api/v1/aiorchestratorownership/commands
GET  /api/v1/aiorchestratorownership/{id}
GET  /api/v1/aiorchestratorownership/{id}/timeline
GET  /api/v1/aiorchestratorownership/{id}/evidence
GET  /api/v1/aiorchestratorownership/{id}/decisions
POST /api/v1/aiorchestratorownership/{id}/exceptions
POST /api/v1/aiorchestratorownership/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aiorchestratorownership_aggregate	authoritative subject
aiorchestratorownership_version	immutable submitted, approved, issued, or signed versions
aiorchestratorownership_relationship	governed relationships
aiorchestratorownership_decision	binding Decisions and rationale
aiorchestratorownership_evidence_link	provenance and evidence relationships
aiorchestratorownership_event_outbox	transactional Event publication
aiorchestratorownership_idempotency	duplicate-write protection
aiorchestratorownership_exception	exception, claim, dispute, or hold
aiorchestratorownership_correction	correction and supersession lineage
aiorchestratorownership_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Agent Run State Machine

Metadata	Value
Volume	9
Chapter ID	V09-04
Subject	AgentRunStateMachine
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Agent Run State Machine.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AgentRunStateMachine	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AgentRunStateMachineVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AgentRunStateMachineDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AgentRunStateMachineEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AgentRunStateMachineException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AgentRunStateMachine {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAgentRunStateMachine
- ValidateAgentRunStateMachine
- ApproveAgentRunStateMachine
- ActivateAgentRunStateMachine
- SuspendAgentRunStateMachine
- CorrectAgentRunStateMachine
- RetireAgentRunStateMachine

Reference Events:

- AgentRunStateMachineCreated
- AgentRunStateMachineValidated
- AgentRunStateMachineApproved
- AgentRunStateMachineActivated
- AgentRunStateMachineSuspended
- AgentRunStateMachineCorrected
- AgentRunStateMachineRetired

Reference queries:

- GetAgentRunStateMachine
- ListAgentRunStateMachineRecords
- GetAgentRunStateMachineTimeline
- GetAgentRunStateMachineEvidence
- GetAgentRunStateMachineDecisionPackage
- GetAgentRunStateMachineExceptions
- GetAgentRunStateMachineAuditTrail

Reference API surface:

```

POST /api/v1/agentrunstatemachine/commands
GET  /api/v1/agentrunstatemachine/{id}
GET  /api/v1/agentrunstatemachine/{id}/timeline
GET  /api/v1/agentrunstatemachine/{id}/evidence
GET  /api/v1/agentrunstatemachine/{id}/decisions
POST /api/v1/agentrunstatemachine/{id}/exceptions
POST /api/v1/agentrunstatemachine/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
agentrunstatemachine_aggregate	authoritative subject
agentrunstatemachine_version	immutable submitted, approved, issued, or signed versions
agentrunstatemachine_relationship	governed relationships
agentrunstatemachine_decision	binding Decisions and rationale
agentrunstatemachine_evidence_link	provenance and evidence relationships
agentrunstatemachine_event_outbox	transactional Event publication
agentrunstatemachine_idempotency	duplicate-write protection
agentrunstatemachine_exception	exception, claim, dispute, or hold
agentrunstatemachine_correction	correction and supersession lineage
agentrunstatemachine_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Supervisor and Delegation

Metadata	Value
Volume	9
Chapter ID	V09-05
Subject	SupervisorandDelegation
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Supervisor and Delegation.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
SupervisorandDelegation	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SupervisorandDelegationVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SupervisorandDelegationDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SupervisorandDelegationEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SupervisorandDelegationException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
SupervisorandDelegation {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```



```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateSupervisorandDelegation
- ValidateSupervisorandDelegation
- ApproveSupervisorandDelegation
- ActivateSupervisorandDelegation
- SuspendSupervisorandDelegation
- CorrectSupervisorandDelegation
- RetireSupervisorandDelegation

Reference Events:

- SupervisorandDelegationCreated
- SupervisorandDelegationValidated
- SupervisorandDelegationApproved
- SupervisorandDelegationActivated
- SupervisorandDelegationSuspended
- SupervisorandDelegationCorrected
- SupervisorandDelegationRetired

Reference queries:

- GetSupervisorandDelegation
- ListSupervisorandDelegationRecords
- GetSupervisorandDelegationTimeline
- GetSupervisorandDelegationEvidence
- GetSupervisorandDelegationDecisionPackage
- GetSupervisorandDelegationExceptions
- GetSupervisorandDelegationAuditTrail

Reference API surface:

```

POST /api/v1/supervisoranddelegation/commands
GET  /api/v1/supervisoranddelegation/{id}
GET  /api/v1/supervisoranddelegation/{id}/timeline
GET  /api/v1/supervisoranddelegation/{id}/evidence
GET  /api/v1/supervisoranddelegation/{id}/decisions
POST /api/v1/supervisoranddelegation/{id}/exceptions
POST /api/v1/supervisoranddelegation/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
supervisoranddelegation_aggregate	authoritative subject
supervisoranddelegation_version	immutable submitted, approved, issued, or signed versions
supervisoranddelegation_relationship	governed relationships
supervisoranddelegation_decision	binding Decisions and rationale
supervisoranddelegation_evidence_link	provenance and evidence relationships
supervisoranddelegation_event_outbox	transactional Event publication
supervisoranddelegation_idempotency	duplicate-write protection
supervisoranddelegation_exception	exception, claim, dispute, or hold
supervisoranddelegation_correction	correction and supersession lineage
supervisoranddelegation_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Context Builder

Metadata	Value
Volume	9
Chapter ID	V09-06
Subject	ContextBuilder
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Context Builder.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
ContextBuilder	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ContextBuilderVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ContextBuilderDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ContextBuilderEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ContextBuilderException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
ContextBuilder {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateContextBuilder
- ValidateContextBuilder
- ApproveContextBuilder
- ActivateContextBuilder
- SuspendContextBuilder
- CorrectContextBuilder
- RetireContextBuilder

Reference Events:

- ContextBuilderCreated
- ContextBuilderValidated
- ContextBuilderApproved
- ContextBuilderActivated
- ContextBuilderSuspended
- ContextBuilderCorrected
- ContextBuilderRetired

Reference queries:

- GetContextBuilder
- ListContextBuilderRecords
- GetContextBuilderTimeline
- GetContextBuilderEvidence
- GetContextBuilderDecisionPackage
- GetContextBuilderExceptions
- GetContextBuilderAuditTrail

Reference API surface:

```

POST /api/v1/contextbuilder/commands
GET  /api/v1/contextbuilder/{id}
GET  /api/v1/contextbuilder/{id}/timeline
GET  /api/v1/contextbuilder/{id}/evidence
GET  /api/v1/contextbuilder/{id}/decisions
POST /api/v1/contextbuilder/{id}/exceptions
POST /api/v1/contextbuilder/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```



```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
contextbuilder_aggregate	authoritative subject
contextbuilder_version	immutable submitted, approved, issued, or signed versions
contextbuilder_relationship	governed relationships
contextbuilder_decision	binding Decisions and rationale
contextbuilder_evidence_link	provenance and evidence relationships
contextbuilder_event_outbox	transactional Event publication
contextbuilder_idempotency	duplicate-write protection
contextbuilder_exception	exception, claim, dispute, or hold
contextbuilder_correction	correction and supersession lineage
contextbuilder_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Planner and Plan Versioning

Metadata	Value
Volume	9
Chapter ID	V09-07
Subject	PlannerandPlanVersioning
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Planner and Plan Versioning**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
PlannerandPlanVersioning	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlannerandPlanVersioningVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlannerandPlanVersioningDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlannerandPlanVersioningEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PlannerandPlanVersioningException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
PlannerandPlanVersioning {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreatePlannerandPlanVersioning
- ValidatePlannerandPlanVersioning
- ApprovePlannerandPlanVersioning
- ActivatePlannerandPlanVersioning
- SuspendPlannerandPlanVersioning
- CorrectPlannerandPlanVersioning
- RetirePlannerandPlanVersioning

Reference Events:

- PlannerandPlanVersioningCreated
- PlannerandPlanVersioningValidated
- PlannerandPlanVersioningApproved
- PlannerandPlanVersioningActivated
- PlannerandPlanVersioningSuspended
- PlannerandPlanVersioningCorrected
- PlannerandPlanVersioningRetired

Reference queries:

- GetPlannerandPlanVersioning
- ListPlannerandPlanVersioningRecords
- GetPlannerandPlanVersioningTimeline
- GetPlannerandPlanVersioningEvidence
- GetPlannerandPlanVersioningDecisionPackage
- GetPlannerandPlanVersioningExceptions
- GetPlannerandPlanVersioningAuditTrail

Reference API surface:

```

POST /api/v1/plannerandplanversioning/commands
GET  /api/v1/plannerandplanversioning/{id}
GET  /api/v1/plannerandplanversioning/{id}/timeline
GET  /api/v1/plannerandplanversioning/{id}/evidence
GET  /api/v1/plannerandplanversioning/{id}/decisions
POST /api/v1/plannerandplanversioning/{id}/exceptions
POST /api/v1/plannerandplanversioning/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
plannerandplanversioning_aggregate	authoritative subject
plannerandplanversioning_version	immutable submitted, approved, issued, or signed versions
plannerandplanversioning_relationship	governed relationships
plannerandplanversioning_decision	binding Decisions and rationale
plannerandplanversioning_evidence_link	provenance and evidence relationships
plannerandplanversioning_event_outbox	transactional Event publication
plannerandplanversioning_idempotency	duplicate-write protection
plannerandplanversioning_exception	exception, claim, dispute, or hold
plannerandplanversioning_correction	correction and supersession lineage
plannerandplanversioning_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Critic and Verifier

Metadata	Value
Volume	9
Chapter ID	V09-08
Subject	CriticandVerifier
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Critic and Verifier.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
CriticandVerifier	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CriticandVerifierVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CriticandVerifierDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CriticandVerifierEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CriticandVerifierException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
CriticandVerifier {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateCriticandVerifier
- ValidateCriticandVerifier
- ApproveCriticandVerifier
- ActivateCriticandVerifier
- SuspendCriticandVerifier
- CorrectCriticandVerifier
- RetireCriticandVerifier

Reference Events:

- CriticandVerifierCreated
- CriticandVerifierValidated
- CriticandVerifierApproved
- CriticandVerifierActivated
- CriticandVerifierSuspended
- CriticandVerifierCorrected
- CriticandVerifierRetired

Reference queries:

- GetCriticandVerifier
- ListCriticandVerifierRecords
- GetCriticandVerifierTimeline
- GetCriticandVerifierEvidence
- GetCriticandVerifierDecisionPackage
- GetCriticandVerifierExceptions
- GetCriticandVerifierAuditTrail

Reference API surface:

```

POST /api/v1/criticandverifier/commands
GET  /api/v1/criticandverifier/{id}
GET  /api/v1/criticandverifier/{id}/timeline
GET  /api/v1/criticandverifier/{id}/evidence
GET  /api/v1/criticandverifier/{id}/decisions
POST /api/v1/criticandverifier/{id}/exceptions
POST /api/v1/criticandverifier/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
criticandverifier_aggregate	authoritative subject
criticandverifier_version	immutable submitted, approved, issued, or signed versions
criticandverifier_relationship	governed relationships
criticandverifier_decision	binding Decisions and rationale
criticandverifier_evidence_link	provenance and evidence relationships
criticandverifier_event_outbox	transactional Event publication
criticandverifier_idempotency	duplicate-write protection
criticandverifier_exception	exception, claim, dispute, or hold
criticandverifier_correction	correction and supersession lineage
criticandverifier_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

Policy and Approval Engine

Metadata	Value
Volume	9
Chapter ID	V09-09
Subject	PolicyandApprovalEngine
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Policy and Approval Engine.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
PolicyandApprovalEngine	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PolicyandApprovalEngineVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PolicyandApprovalEngineDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PolicyandApprovalEngineEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PolicyandApprovalEngineException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
PolicyandApprovalEngine {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreatePolicyandApprovalEngine
- ValidatePolicyandApprovalEngine
- ApprovePolicyandApprovalEngine
- ActivatePolicyandApprovalEngine
- SuspendPolicyandApprovalEngine
- CorrectPolicyandApprovalEngine
- RetirePolicyandApprovalEngine

Reference Events:

- PolicyandApprovalEngineCreated
- PolicyandApprovalEngineValidated
- PolicyandApprovalEngineApproved
- PolicyandApprovalEngineActivated
- PolicyandApprovalEngineSuspended
- PolicyandApprovalEngineCorrected
- PolicyandApprovalEngineRetired

Reference queries:

- GetPolicyandApprovalEngine
- ListPolicyandApprovalEngineRecords
- GetPolicyandApprovalEngineTimeline
- GetPolicyandApprovalEngineEvidence
- GetPolicyandApprovalEngineDecisionPackage
- GetPolicyandApprovalEngineExceptions
- GetPolicyandApprovalEngineAuditTrail

Reference API surface:

```

POST /api/v1/policyandapprovalengine/commands
GET  /api/v1/policyandapprovalengine/{id}
GET  /api/v1/policyandapprovalengine/{id}/timeline
GET  /api/v1/policyandapprovalengine/{id}/evidence
GET  /api/v1/policyandapprovalengine/{id}/decisions
POST /api/v1/policyandapprovalengine/{id}/exceptions
POST /api/v1/policyandapprovalengine/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
policyandapprovalengine_aggregate	authoritative subject
policyandapprovalengine_version	immutable submitted, approved, issued, or signed versions
policyandapprovalengine_relationship	governed relationships
policyandapprovalengine_decision	binding Decisions and rationale
policyandapprovalengine_evidence_link	provenance and evidence relationships
policyandapprovalengine_event_outbox	transactional Event publication
policyandapprovalengine_idempotency	duplicate-write protection
policyandapprovalengine_exception	exception, claim, dispute, or hold
policyandapprovalengine_correction	correction and supersession lineage
policyandapprovalengine_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Typed Tool Executor

Metadata	Value
Volume	9
Chapter ID	V09-10
Subject	TypedToolExecutor
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Typed Tool Executor**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
TypedToolExecutor	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
TypedToolExecutorVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
TypedToolExecutorDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
TypedToolExecutorEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
TypedToolExecutorException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
TypedToolExecutor {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateTypedToolExecutor
- ValidateTypedToolExecutor
- ApproveTypedToolExecutor
- ActivateTypedToolExecutor
- SuspendTypedToolExecutor
- CorrectTypedToolExecutor
- RetireTypedToolExecutor

Reference Events:

- TypedToolExecutorCreated
- TypedToolExecutorValidated
- TypedToolExecutorApproved
- TypedToolExecutorActivated
- TypedToolExecutorSuspended
- TypedToolExecutorCorrected
- TypedToolExecutorRetired

Reference queries:

- GetTypedToolExecutor
- ListTypedToolExecutorRecords
- GetTypedToolExecutorTimeline
- GetTypedToolExecutorEvidence
- GetTypedToolExecutorDecisionPackage
- GetTypedToolExecutorExceptions
- GetTypedToolExecutorAuditTrail

Reference API surface:

```

POST /api/v1/typedtoolexecutor/commands
GET  /api/v1/typedtoolexecutor/{id}
GET  /api/v1/typedtoolexecutor/{id}/timeline
GET  /api/v1/typedtoolexecutor/{id}/evidence
GET  /api/v1/typedtoolexecutor/{id}/decisions
POST /api/v1/typedtoolexecutor/{id}/exceptions
POST /api/v1/typedtoolexecutor/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
typedtoolexecutor_aggregate	authoritative subject
typedtoolexecutor_version	immutable submitted, approved, issued, or signed versions
typedtoolexecutor_relationship	governed relationships
typedtoolexecutor_decision	binding Decisions and rationale
typedtoolexecutor_evidence_link	provenance and evidence relationships
typedtoolexecutor_event_outbox	transactional Event publication
typedtoolexecutor_idempotency	duplicate-write protection
typedtoolexecutor_exception	exception, claim, dispute, or hold
typedtoolexecutor_correction	correction and supersession lineage
typedtoolexecutor_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Outcome Observer

Metadata	Value
Volume	9
Chapter ID	V09-11
Subject	OutcomeObserver
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Outcome Observer.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
OutcomeObserver	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OutcomeObserverVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OutcomeObserverDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OutcomeObserverEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OutcomeObserverException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
OutcomeObserver {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateOutcomeObserver
- ValidateOutcomeObserver
- ApproveOutcomeObserver
- ActivateOutcomeObserver
- SuspendOutcomeObserver
- CorrectOutcomeObserver
- RetireOutcomeObserver

Reference Events:

- OutcomeObserverCreated
- OutcomeObserverValidated
- OutcomeObserverApproved
- OutcomeObserverActivated
- OutcomeObserverSuspended
- OutcomeObserverCorrected
- OutcomeObserverRetired

Reference queries:

- GetOutcomeObserver
- ListOutcomeObserverRecords
- GetOutcomeObserverTimeline
- GetOutcomeObserverEvidence
- GetOutcomeObserverDecisionPackage
- GetOutcomeObserverExceptions
- GetOutcomeObserverAuditTrail

Reference API surface:

```

POST /api/v1/outcomeobserver/commands
GET  /api/v1/outcomeobserver/{id}
GET  /api/v1/outcomeobserver/{id}/timeline
GET  /api/v1/outcomeobserver/{id}/evidence
GET  /api/v1/outcomeobserver/{id}/decisions
POST /api/v1/outcomeobserver/{id}/exceptions
POST /api/v1/outcomeobserver/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
outcomeobserver_aggregate	authoritative subject
outcomeobserver_version	immutable submitted, approved, issued, or signed versions
outcomeobserver_relationship	governed relationships
outcomeobserver_decision	binding Decisions and rationale
outcomeobserver_evidence_link	provenance and evidence relationships
outcomeobserver_event_outbox	transactional Event publication
outcomeobserver_idempotency	duplicate-write protection
outcomeobserver_exception	exception, claim, dispute, or hold
outcomeobserver_correction	correction and supersession lineage
outcomeobserver_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

Recovery and Replanner

Metadata	Value
Volume	9
Chapter ID	V09-12
Subject	RecoveryandReplanner
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Recovery and Replanner.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
RecoveryandReplanner	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RecoveryandReplannerVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RecoveryandReplannerDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RecoveryandReplannerEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RecoveryandReplannerException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
RecoveryandReplanner {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateRecoveryandReplanner
- ValidateRecoveryandReplanner
- ApproveRecoveryandReplanner
- ActivateRecoveryandReplanner
- SuspendRecoveryandReplanner
- CorrectRecoveryandReplanner
- RetireRecoveryandReplanner

Reference Events:

- RecoveryandReplannerCreated
- RecoveryandReplannerValidated
- RecoveryandReplannerApproved
- RecoveryandReplannerActivated
- RecoveryandReplannerSuspended
- RecoveryandReplannerCorrected
- RecoveryandReplannerRetired

Reference queries:

- GetRecoveryandReplanner
- ListRecoveryandReplannerRecords
- GetRecoveryandReplannerTimeline
- GetRecoveryandReplannerEvidence
- GetRecoveryandReplannerDecisionPackage
- GetRecoveryandReplannerExceptions
- GetRecoveryandReplannerAuditTrail

Reference API surface:

```

POST /api/v1/recoveryandreplanner/commands
GET  /api/v1/recoveryandreplanner/{id}
GET  /api/v1/recoveryandreplanner/{id}/timeline
GET  /api/v1/recoveryandreplanner/{id}/evidence
GET  /api/v1/recoveryandreplanner/{id}/decisions
POST /api/v1/recoveryandreplanner/{id}/exceptions
POST /api/v1/recoveryandreplanner/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
recoveryandreplanner_aggregate	authoritative subject
recoveryandreplanner_version	immutable submitted, approved, issued, or signed versions
recoveryandreplanner_relationship	governed relationships
recoveryandreplanner_decision	binding Decisions and rationale
recoveryandreplanner_evidence_link	provenance and evidence relationships
recoveryandreplanner_event_outbox	transactional Event publication
recoveryandreplanner_idempotency	duplicate-write protection
recoveryandreplanner_exception	exception, claim, dispute, or hold
recoveryandreplanner_correction	correction and supersession lineage
recoveryandreplanner_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

MCP Tool Registry

Metadata	Value
Volume	9
Chapter ID	V09-13
Subject	MCPToolRegistry
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for MCP Tool Registry**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
MCPToolRegistry	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolRegistryVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolRegistryDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolRegistryEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolRegistryException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
MCPToolRegistry {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```



```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateMCPToolRegistry
- ValidateMCPToolRegistry
- ApproveMCPToolRegistry
- ActivateMCPToolRegistry
- SuspendMCPToolRegistry
- CorrectMCPToolRegistry
- RetireMCPToolRegistry

Reference Events:

- MCPToolRegistryCreated
- MCPToolRegistryValidated
- MCPToolRegistryApproved
- MCPToolRegistryActivated
- MCPToolRegistrySuspended
- MCPToolRegistryCorrected
- MCPToolRegistryRetired

Reference queries:

- GetMCPToolRegistry
- ListMCPToolRegistryRecords
- GetMCPToolRegistryTimeline
- GetMCPToolRegistryEvidence
- GetMCPToolRegistryDecisionPackage
- GetMCPToolRegistryExceptions
- GetMCPToolRegistryAuditTrail

Reference API surface:

```

POST /api/v1/mcptoolregistry/commands
GET  /api/v1/mcptoolregistry/{id}
GET  /api/v1/mcptoolregistry/{id}/timeline
GET  /api/v1/mcptoolregistry/{id}/evidence
GET  /api/v1/mcptoolregistry/{id}/decisions
POST /api/v1/mcptoolregistry/{id}/exceptions
POST /api/v1/mcptoolregistry/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
mcptoolregistry_aggregate	authoritative subject
mcptoolregistry_version	immutable submitted, approved, issued, or signed versions
mcptoolregistry_relationship	governed relationships
mcptoolregistry_decision	binding Decisions and rationale
mcptoolregistry_evidence_link	provenance and evidence relationships
mcptoolregistry_event_outbox	transactional Event publication
mcptoolregistry_idempotency	duplicate-write protection
mcptoolregistry_exception	exception, claim, dispute, or hold
mcptoolregistry_correction	correction and supersession lineage
mcptoolregistry_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

MCP Tool Gateway

Metadata	Value
Volume	9
Chapter ID	V09-14
Subject	MCPToolGateway
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for MCP Tool Gateway**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
MCPToolGateway	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolGatewayVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolGatewayDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolGatewayEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
MCPToolGatewayException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
MCPToolGateway {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateMCPToolGateway
- ValidateMCPToolGateway
- ApproveMCPToolGateway
- ActivateMCPToolGateway
- SuspendMCPToolGateway
- CorrectMCPToolGateway
- RetireMCPToolGateway

Reference Events:

- MCPToolGatewayCreated
- MCPToolGatewayValidated
- MCPToolGatewayApproved
- MCPToolGatewayActivated
- MCPToolGatewaySuspended
- MCPToolGatewayCorrected
- MCPToolGatewayRetired

Reference queries:

- GetMCPToolGateway
- ListMCPToolGatewayRecords
- GetMCPToolGatewayTimeline
- GetMCPToolGatewayEvidence
- GetMCPToolGatewayDecisionPackage
- GetMCPToolGatewayExceptions
- GetMCPToolGatewayAuditTrail

Reference API surface:

```

POST /api/v1/mcptoolgateway/commands
GET  /api/v1/mcptoolgateway/{id}
GET  /api/v1/mcptoolgateway/{id}/timeline
GET  /api/v1/mcptoolgateway/{id}/evidence
GET  /api/v1/mcptoolgateway/{id}/decisions
POST /api/v1/mcptoolgateway/{id}/exceptions
POST /api/v1/mcptoolgateway/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```



```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
mcptoolgateway_aggregate	authoritative subject
mcptoolgateway_version	immutable submitted, approved, issued, or signed versions
mcptoolgateway_relationship	governed relationships
mcptoolgateway_decision	binding Decisions and rationale
mcptoolgateway_evidence_link	provenance and evidence relationships
mcptoolgateway_event_outbox	transactional Event publication
mcptoolgateway_idempotency	duplicate-write protection
mcptoolgateway_exception	exception, claim, dispute, or hold
mcptoolgateway_correction	correction and supersession lineage
mcptoolgateway_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Workflow Engine Integration

Metadata	Value
Volume	9
Chapter ID	V09-15
Subject	WorkflowEngineIntegration
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Workflow Engine Integration.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
WorkflowEngineIntegration	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
WorkflowEngineIntegrationVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
WorkflowEngineIntegrationDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
WorkflowEngineIntegrationEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
WorkflowEngineIntegrationException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
WorkflowEngineIntegration {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateWorkflowEngineIntegration
- ValidateWorkflowEngineIntegration
- ApproveWorkflowEngineIntegration
- ActivateWorkflowEngineIntegration
- SuspendWorkflowEngineIntegration
- CorrectWorkflowEngineIntegration
- RetireWorkflowEngineIntegration

Reference Events:

- WorkflowEngineIntegrationCreated
- WorkflowEngineIntegrationValidated
- WorkflowEngineIntegrationApproved
- WorkflowEngineIntegrationActivated
- WorkflowEngineIntegrationSuspended
- WorkflowEngineIntegrationCorrected
- WorkflowEngineIntegrationRetired

Reference queries:

- GetWorkflowEngineIntegration
- ListWorkflowEngineIntegrationRecords
- GetWorkflowEngineIntegrationTimeline
- GetWorkflowEngineIntegrationEvidence
- GetWorkflowEngineIntegrationDecisionPackage
- GetWorkflowEngineIntegrationExceptions
- GetWorkflowEngineIntegrationAuditTrail

Reference API surface:

```

POST /api/v1/workflowengineintegration/commands
GET  /api/v1/workflowengineintegration/{id}
GET  /api/v1/workflowengineintegration/{id}/timeline
GET  /api/v1/workflowengineintegration/{id}/evidence
GET  /api/v1/workflowengineintegration/{id}/decisions
POST /api/v1/workflowengineintegration/{id}/exceptions
POST /api/v1/workflowengineintegration/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
workflowengineintegration_aggregate	authoritative subject
workflowengineintegration_version	immutable submitted, approved, issued, or signed versions
workflowengineintegration_relationship	governed relationships
workflowengineintegration_decision	binding Decisions and rationale
workflowengineintegration_evidence_link	provenance and evidence relationships
workflowengineintegration_event_outbox	transactional Event publication
workflowengineintegration_idempotency	duplicate-write protection
workflowengineintegration_exception	exception, claim, dispute, or hold
workflowengineintegration_correction	correction and supersession lineage
workflowengineintegration_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Event Bus Integration

Metadata	Value
Volume	9
Chapter ID	V09-16
Subject	EventBusIntegration
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Event Bus Integration.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
EventBusIntegration	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EventBusIntegrationVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EventBusIntegrationDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EventBusIntegrationEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EventBusIntegrationException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

EventBusIntegration {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateEventBusIntegration
- ValidateEventBusIntegration
- ApproveEventBusIntegration
- ActivateEventBusIntegration
- SuspendEventBusIntegration
- CorrectEventBusIntegration
- RetireEventBusIntegration

Reference Events:

- EventBusIntegrationCreated
- EventBusIntegrationValidated
- EventBusIntegrationApproved
- EventBusIntegrationActivated
- EventBusIntegrationSuspended
- EventBusIntegrationCorrected
- EventBusIntegrationRetired

Reference queries:

- GetEventBusIntegration
- ListEventBusIntegrationRecords
- GetEventBusIntegrationTimeline
- GetEventBusIntegrationEvidence
- GetEventBusIntegrationDecisionPackage
- GetEventBusIntegrationExceptions
- GetEventBusIntegrationAuditTrail

Reference API surface:

```

POST /api/v1/eventbusintegration/commands
GET  /api/v1/eventbusintegration/{id}
GET  /api/v1/eventbusintegration/{id}/timeline
GET  /api/v1/eventbusintegration/{id}/evidence
GET  /api/v1/eventbusintegration/{id}/decisions
POST /api/v1/eventbusintegration/{id}/exceptions
POST /api/v1/eventbusintegration/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
eventbusintegration_aggregate	authoritative subject
eventbusintegration_version	immutable submitted, approved, issued, or signed versions
eventbusintegration_relationship	governed relationships
eventbusintegration_decision	binding Decisions and rationale
eventbusintegration_evidence_link	provenance and evidence relationships
eventbusintegration_event_outbox	transactional Event publication
eventbusintegration_idempotency	duplicate-write protection
eventbusintegration_exception	exception, claim, dispute, or hold
eventbusintegration_correction	correction and supersession lineage
eventbusintegration_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Provider Adapter Architecture

Metadata	Value
Volume	9
Chapter ID	V09-17
Subject	ProviderAdapterArchitecture
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Provider Adapter Architecture**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
ProviderAdapterArchitecture	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ProviderAdapterArchitectureVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ProviderAdapterArchitectureDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ProviderAdapterArchitectureEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ProviderAdapterArchitectureException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
ProviderAdapterArchitecture {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateProviderAdapterArchitecture
- ValidateProviderAdapterArchitecture
- ApproveProviderAdapterArchitecture
- ActivateProviderAdapterArchitecture
- SuspendProviderAdapterArchitecture
- CorrectProviderAdapterArchitecture
- RetireProviderAdapterArchitecture

Reference Events:

- ProviderAdapterArchitectureCreated
- ProviderAdapterArchitectureValidated
- ProviderAdapterArchitectureApproved
- ProviderAdapterArchitectureActivated
- ProviderAdapterArchitectureSuspended
- ProviderAdapterArchitectureCorrected
- ProviderAdapterArchitectureRetired

Reference queries:

- GetProviderAdapterArchitecture
- ListProviderAdapterArchitectureRecords
- GetProviderAdapterArchitectureTimeline
- GetProviderAdapterArchitectureEvidence
- GetProviderAdapterArchitectureDecisionPackage
- GetProviderAdapterArchitectureExceptions
- GetProviderAdapterArchitectureAuditTrail

Reference API surface:

```

POST /api/v1/provideradapterarchitecture/commands
GET  /api/v1/provideradapterarchitecture/{id}
GET  /api/v1/provideradapterarchitecture/{id}/timeline
GET  /api/v1/provideradapterarchitecture/{id}/evidence
GET  /api/v1/provideradapterarchitecture/{id}/decisions
POST /api/v1/provideradapterarchitecture/{id}/exceptions
POST /api/v1/provideradapterarchitecture/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
provideradapterarchitecture_aggregate	authoritative subject
provideradapterarchitecture_version	immutable submitted, approved, issued, or signed versions
provideradapterarchitecture_relationship	governed relationships
provideradapterarchitecture_decision	binding Decisions and rationale
provideradapterarchitecture_evidence_link	provenance and evidence relationships
provideradapterarchitecture_event_outbox	transactional Event publication
provideradapterarchitecture_idempotency	duplicate-write protection
provideradapterarchitecture_exception	exception, claim, dispute, or hold
provideradapterarchitecture_correction	correction and supersession lineage
provideradapterarchitecture_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Model Registry and Routing

Metadata	Value
Volume	9
Chapter ID	V09-18
Subject	ModelRegistryandRouting
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Model Registry and Routing.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
ModelRegistryandRouting	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ModelRegistryandRoutingVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ModelRegistryandRoutingDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ModelRegistryandRoutingEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
ModelRegistryandRoutingException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
ModelRegistryandRouting {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateModelRegistryandRouting
- ValidateModelRegistryandRouting
- ApproveModelRegistryandRouting
- ActivateModelRegistryandRouting
- SuspendModelRegistryandRouting
- CorrectModelRegistryandRouting
- RetireModelRegistryandRouting

Reference Events:

- ModelRegistryandRoutingCreated
- ModelRegistryandRoutingValidated
- ModelRegistryandRoutingApproved
- ModelRegistryandRoutingActivated
- ModelRegistryandRoutingSuspended
- ModelRegistryandRoutingCorrected
- ModelRegistryandRoutingRetired

Reference queries:

- GetModelRegistryandRouting
- ListModelRegistryandRoutingRecords
- GetModelRegistryandRoutingTimeline
- GetModelRegistryandRoutingEvidence
- GetModelRegistryandRoutingDecisionPackage
- GetModelRegistryandRoutingExceptions
- GetModelRegistryandRoutingAuditTrail

Reference API surface:

```

POST /api/v1/modelregistryandrouting/commands
GET  /api/v1/modelregistryandrouting/{id}
GET  /api/v1/modelregistryandrouting/{id}/timeline
GET  /api/v1/modelregistryandrouting/{id}/evidence
GET  /api/v1/modelregistryandrouting/{id}/decisions
POST /api/v1/modelregistryandrouting/{id}/exceptions
POST /api/v1/modelregistryandrouting/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
modelregistryandrouting_aggregate	authoritative subject
modelregistryandrouting_version	immutable submitted, approved, issued, or signed versions
modelregistryandrouting_relationship	governed relationships
modelregistryandrouting_decision	binding Decisions and rationale
modelregistryandrouting_evidence_link	provenance and evidence relationships
modelregistryandrouting_event_outbox	transactional Event publication
modelregistryandrouting_idempotency	duplicate-write protection
modelregistryandrouting_exception	exception, claim, dispute, or hold
modelregistryandrouting_correction	correction and supersession lineage
modelregistryandrouting_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Prompt and Template Governance

Metadata	Value
Volume	9
Chapter ID	V09-19
Subject	PromptandTemplateGovernance
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Prompt and Template Governance**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
PromptandTemplateGovernance	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PromptandTemplateGovernanceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PromptandTemplateGovernanceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PromptandTemplateGovernanceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PromptandTemplateGovernanceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
PromptandTemplateGovernance {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreatePromptandTemplateGovernance
- ValidatePromptandTemplateGovernance
- ApprovePromptandTemplateGovernance
- ActivatePromptandTemplateGovernance
- SuspendPromptandTemplateGovernance
- CorrectPromptandTemplateGovernance
- RetirePromptandTemplateGovernance

Reference Events:

- PromptandTemplateGovernanceCreated
- PromptandTemplateGovernanceValidated
- PromptandTemplateGovernanceApproved
- PromptandTemplateGovernanceActivated
- PromptandTemplateGovernanceSuspended
- PromptandTemplateGovernanceCorrected
- PromptandTemplateGovernanceRetired

Reference queries:

- GetPromptandTemplateGovernance
- ListPromptandTemplateGovernanceRecords
- GetPromptandTemplateGovernanceTimeline
- GetPromptandTemplateGovernanceEvidence
- GetPromptandTemplateGovernanceDecisionPackage
- GetPromptandTemplateGovernanceExceptions
- GetPromptandTemplateGovernanceAuditTrail

Reference API surface:

```

POST /api/v1/promptandtemplategovernance/commands
GET  /api/v1/promptandtemplategovernance/{id}
GET  /api/v1/promptandtemplategovernance/{id}/timeline
GET  /api/v1/promptandtemplategovernance/{id}/evidence
GET  /api/v1/promptandtemplategovernance/{id}/decisions
POST /api/v1/promptandtemplategovernance/{id}/exceptions
POST /api/v1/promptandtemplategovernance/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
promptandtemplategovernance_aggregate	authoritative subject
promptandtemplategovernance_version	immutable submitted, approved, issued, or signed versions
promptandtemplategovernance_relationship	governed relationships
promptandtemplategovernance_decision	binding Decisions and rationale
promptandtemplategovernance_evidence_link	provenance and evidence relationships
promptandtemplategovernance_event_outbox	transactional Event publication
promptandtemplategovernance_idempotency	duplicate-write protection
promptandtemplategovernance_exception	exception, claim, dispute, or hold
promptandtemplategovernance_correction	correction and supersession lineage
promptandtemplategovernance_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Structured Output and Schema Validation

Metadata	Value
Volume	9
Chapter ID	V09-20
Subject	StructuredOutputandSchemaValidation
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Structured Output and Schema Validation**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome
≠ Observation
≠ Claim
≠ derived Perspective
≠ Prediction
≠ Recommendation
≠ authorized Decision
≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;

- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
StructuredOutputandSchemaValidation	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
StructuredOutputandSchemaValidationVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
StructuredOutputandSchemaValidationDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
StructuredOutputandSchemaValidationEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
StructuredOutputandSchemaValidationException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
StructuredOutputandSchemaValidation {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage

- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
```

```

decisionType
subjectRef
sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateStructuredOutputandSchemaValidation
- ValidateStructuredOutputandSchemaValidation
- ApproveStructuredOutputandSchemaValidation
- ActivateStructuredOutputandSchemaValidation
- SuspendStructuredOutputandSchemaValidation
- CorrectStructuredOutputandSchemaValidation
- RetireStructuredOutputandSchemaValidation

Reference Events:

- StructuredOutputandSchemaValidationCreated
- StructuredOutputandSchemaValidationValidated
- StructuredOutputandSchemaValidationApproved
- StructuredOutputandSchemaValidationActivated
- StructuredOutputandSchemaValidationSuspended
- StructuredOutputandSchemaValidationCorrected
- StructuredOutputandSchemaValidationRetired

Reference queries:

- GetStructuredOutputandSchemaValidation
- ListStructuredOutputandSchemaValidationRecords
- GetStructuredOutputandSchemaValidationTimeline
- GetStructuredOutputandSchemaValidationEvidence
- GetStructuredOutputandSchemaValidationDecisionPackage
- GetStructuredOutputandSchemaValidationExceptions
- GetStructuredOutputandSchemaValidationAuditTrail

Reference API surface:

```

POST /api/v1/structuredoutputandschemavalidation/commands
GET  /api/v1/structuredoutputandschemavalidation/{id}
GET  /api/v1/structuredoutputandschemavalidation/{id}/timeline
GET  /api/v1/structuredoutputandschemavalidation/{id}/evidence
GET  /api/v1/structuredoutputandschemavalidation/{id}/decisions
POST /api/v1/structuredoutputandschemavalidation/{id}/exceptions
POST /api/v1/structuredoutputandschemavalidation/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B

```

```

C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
structuredoutputandschemavalidation_aggregate	authoritative subject
structuredoutputandschemavalidation_version	immutable submitted, approved, issued, or signed versions
structuredoutputandschemavalidation_relationship	governed relationships
structuredoutputandschemavalidation_decision	binding Decisions and rationale
structuredoutputandschemavalidation_evidence_link	provenance and evidence relationships
structuredoutputandschemavalidation_event_outbox	transactional Event publication
structuredoutputandschemavalidation_idempotency	duplicate-write protection
structuredoutputandschemavalidation_exception	exception, claim, dispute, or hold
structuredoutputandschemavalidation_correction	correction and supersession lineage
structuredoutputandschemavalidation_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result

- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Memory

Metadata	Value
Volume	9
Chapter ID	V09-21
Subject	AIMemory
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Memory**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIMemory	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMemoryVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMemoryDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMemoryEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMemoryException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIMemory {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```



```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIMemory
- ValidateAIMemory
- ApproveAIMemory
- ActivateAIMemory
- SuspendAIMemory
- CorrectAIMemory
- RetireAIMemory

Reference Events:

- AIMemoryCreated
- AIMemoryValidated
- AIMemoryApproved
- AIMemoryActivated
- AIMemorySuspended
- AIMemoryCorrected
- AIMemoryRetired

Reference queries:

- GetAIMemory
- ListAIMemoryRecords
- GetAIMemoryTimeline
- GetAIMemoryEvidence
- GetAIMemoryDecisionPackage
- GetAIMemoryExceptions
- GetAIMemoryAuditTrail

Reference API surface:

```

POST /api/v1/aimemory/commands
GET  /api/v1/aimemory/{id}
GET  /api/v1/aimemory/{id}/timeline
GET  /api/v1/aimemory/{id}/evidence
GET  /api/v1/aimemory/{id}/decisions
POST /api/v1/aimemory/{id}/exceptions
POST /api/v1/aimemory/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aimemory_aggregate	authoritative subject
aimemory_version	immutable submitted, approved, issued, or signed versions
aimemory_relationship	governed relationships
aimemory_decision	binding Decisions and rationale
aimemory_evidence_link	provenance and evidence relationships
aimemory_event_outbox	transactional Event publication
aimemory_idempotency	duplicate-write protection
aimemory_exception	exception, claim, dispute, or hold
aimemory_correction	correction and supersession lineage
aimemory_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

RAG and Retrieval

Metadata	Value
Volume	9
Chapter ID	V09-22
Subject	RAGandRetrieval
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for RAG and Retrieval.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
RAGandRetrieval	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RAGandRetrievalVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RAGandRetrievalDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RAGandRetrievalEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
RAGandRetrievalException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
RAGandRetrieval {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateRAGandRetrieval
- ValidateRAGandRetrieval
- ApproveRAGandRetrieval
- ActivateRAGandRetrieval
- SuspendRAGandRetrieval
- CorrectRAGandRetrieval
- RetireRAGandRetrieval

Reference Events:

- RAGandRetrievalCreated
- RAGandRetrievalValidated
- RAGandRetrievalApproved
- RAGandRetrievalActivated
- RAGandRetrievalSuspended
- RAGandRetrievalCorrected
- RAGandRetrievalRetired

Reference queries:

- GetRAGandRetrieval
- ListRAGandRetrievalRecords
- GetRAGandRetrievalTimeline
- GetRAGandRetrievalEvidence
- GetRAGandRetrievalDecisionPackage
- GetRAGandRetrievalExceptions
- GetRAGandRetrievalAuditTrail

Reference API surface:

```

POST /api/v1/ragandretrieval/commands
GET  /api/v1/ragandretrieval/{id}
GET  /api/v1/ragandretrieval/{id}/timeline
GET  /api/v1/ragandretrieval/{id}/evidence
GET  /api/v1/ragandretrieval/{id}/decisions
POST /api/v1/ragandretrieval/{id}/exceptions
POST /api/v1/ragandretrieval/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```



```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
ragandretrieval_aggregate	authoritative subject
ragandretrieval_version	immutable submitted, approved, issued, or signed versions
ragandretrieval_relationship	governed relationships
ragandretrieval_decision	binding Decisions and rationale
ragandretrieval_evidence_link	provenance and evidence relationships
ragandretrieval_event_outbox	transactional Event publication
ragandretrieval_idempotency	duplicate-write protection
ragandretrieval_exception	exception, claim, dispute, or hold
ragandretrieval_correction	correction and supersession lineage
ragandretrieval_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Knowledge Cutoff and Provenance

Metadata	Value
Volume	9
Chapter ID	V09-23
Subject	KnowledgeCutoffandProvenance
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Knowledge Cutoff and Provenance.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
KnowledgeCutoffandProvenance	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KnowledgeCutoffandProvenanceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KnowledgeCutoffandProvenanceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KnowledgeCutoffandProvenanceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KnowledgeCutoffandProvenanceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
KnowledgeCutoffandProvenance {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateKnowledgeCutoffandProvenance
- ValidateKnowledgeCutoffandProvenance
- ApproveKnowledgeCutoffandProvenance
- ActivateKnowledgeCutoffandProvenance
- SuspendKnowledgeCutoffandProvenance
- CorrectKnowledgeCutoffandProvenance
- RetireKnowledgeCutoffandProvenance

Reference Events:

- KnowledgeCutoffandProvenanceCreated
- KnowledgeCutoffandProvenanceValidated
- KnowledgeCutoffandProvenanceApproved
- KnowledgeCutoffandProvenanceActivated
- KnowledgeCutoffandProvenanceSuspended
- KnowledgeCutoffandProvenanceCorrected
- KnowledgeCutoffandProvenanceRetired

Reference queries:

- GetKnowledgeCutoffandProvenance
- ListKnowledgeCutoffandProvenanceRecords
- GetKnowledgeCutoffandProvenanceTimeline
- GetKnowledgeCutoffandProvenanceEvidence
- GetKnowledgeCutoffandProvenanceDecisionPackage
- GetKnowledgeCutoffandProvenanceExceptions
- GetKnowledgeCutoffandProvenanceAuditTrail

Reference API surface:

```

POST /api/v1/knowledgecutoffandprovenance/commands
GET  /api/v1/knowledgecutoffandprovenance/{id}
GET  /api/v1/knowledgecutoffandprovenance/{id}/timeline
GET  /api/v1/knowledgecutoffandprovenance/{id}/evidence
GET  /api/v1/knowledgecutoffandprovenance/{id}/decisions
POST /api/v1/knowledgecutoffandprovenance/{id}/exceptions
POST /api/v1/knowledgecutoffandprovenance/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
knowledgecutoffandprovenance_aggregate	authoritative subject
knowledgecutoffandprovenance_version	immutable submitted, approved, issued, or signed versions
knowledgecutoffandprovenance_relationship	governed relationships
knowledgecutoffandprovenance_decision	binding Decisions and rationale
knowledgecutoffandprovenance_evidence_link	provenance and evidence relationships
knowledgecutoffandprovenance_event_outbox	transactional Event publication
knowledgecutoffandprovenance_idempotency	duplicate-write protection
knowledgecutoffandprovenance_exception	exception, claim, dispute, or hold
knowledgecutoffandprovenance_correction	correction and supersession lineage
knowledgecutoffandprovenance_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Digital Twin and World Model

Metadata	Value
Volume	9
Chapter ID	V09-24
Subject	DigitalTwinandWorldModel
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Digital Twin and World Model**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
DigitalTwinandWorldModel	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DigitalTwinandWorldModelVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DigitalTwinandWorldModelDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DigitalTwinandWorldModelEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DigitalTwinandWorldModelException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

DigitalTwinandWorldModel {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateDigitalTwinandWorldModel
- ValidateDigitalTwinandWorldModel
- ApproveDigitalTwinandWorldModel
- ActivateDigitalTwinandWorldModel
- SuspendDigitalTwinandWorldModel
- CorrectDigitalTwinandWorldModel
- RetireDigitalTwinandWorldModel

Reference Events:

- DigitalTwinandWorldModelCreated
- DigitalTwinandWorldModelValidated
- DigitalTwinandWorldModelApproved
- DigitalTwinandWorldModelActivated
- DigitalTwinandWorldModelSuspended
- DigitalTwinandWorldModelCorrected
- DigitalTwinandWorldModelRetired

Reference queries:

- GetDigitalTwinandWorldModel
- ListDigitalTwinandWorldModelRecords
- GetDigitalTwinandWorldModelTimeline
- GetDigitalTwinandWorldModelEvidence
- GetDigitalTwinandWorldModelDecisionPackage
- GetDigitalTwinandWorldModelExceptions
- GetDigitalTwinandWorldModelAuditTrail

Reference API surface:

```

POST /api/v1/digitaltwinandworldmodel/commands
GET  /api/v1/digitaltwinandworldmodel/{id}
GET  /api/v1/digitaltwinandworldmodel/{id}/timeline
GET  /api/v1/digitaltwinandworldmodel/{id}/evidence
GET  /api/v1/digitaltwinandworldmodel/{id}/decisions
POST /api/v1/digitaltwinandworldmodel/{id}/exceptions
POST /api/v1/digitaltwinandworldmodel/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
digitaltwinandworldmodel_aggregate	authoritative subject
digitaltwinandworldmodel_version	immutable submitted, approved, issued, or signed versions
digitaltwinandworldmodel_relationship	governed relationships
digitaltwinandworldmodel_decision	binding Decisions and rationale
digitaltwinandworldmodel_evidence_link	provenance and evidence relationships
digitaltwinandworldmodel_event_outbox	transactional Event publication
digitaltwinandworldmodel_idempotency	duplicate-write protection
digitaltwinandworldmodel_exception	exception, claim, dispute, or hold
digitaltwinandworldmodel_correction	correction and supersession lineage
digitaltwinandworldmodel_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Evaluation and Promotion

Metadata	Value
Volume	9
Chapter ID	V09-25
Subject	EvaluationandPromotion
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Evaluation and Promotion.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
EvaluationandPromotion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EvaluationandPromotionVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EvaluationandPromotionDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EvaluationandPromotionEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
EvaluationandPromotionException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

EvaluationandPromotion {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateEvaluationandPromotion
- ValidateEvaluationandPromotion
- ApproveEvaluationandPromotion
- ActivateEvaluationandPromotion
- SuspendEvaluationandPromotion
- CorrectEvaluationandPromotion
- RetireEvaluationandPromotion

Reference Events:

- EvaluationandPromotionCreated
- EvaluationandPromotionValidated
- EvaluationandPromotionApproved
- EvaluationandPromotionActivated
- EvaluationandPromotionSuspended
- EvaluationandPromotionCorrected
- EvaluationandPromotionRetired

Reference queries:

- GetEvaluationandPromotion
- ListEvaluationandPromotionRecords
- GetEvaluationandPromotionTimeline
- GetEvaluationandPromotionEvidence
- GetEvaluationandPromotionDecisionPackage
- GetEvaluationandPromotionExceptions
- GetEvaluationandPromotionAuditTrail

Reference API surface:

```

POST /api/v1/evaluationandpromotion/commands
GET  /api/v1/evaluationandpromotion/{id}
GET  /api/v1/evaluationandpromotion/{id}/timeline
GET  /api/v1/evaluationandpromotion/{id}/evidence
GET  /api/v1/evaluationandpromotion/{id}/decisions
POST /api/v1/evaluationandpromotion/{id}/exceptions
POST /api/v1/evaluationandpromotion/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
evaluationandpromotion_aggregate	authoritative subject
evaluationandpromotion_version	immutable submitted, approved, issued, or signed versions
evaluationandpromotion_relationship	governed relationships
evaluationandpromotion_decision	binding Decisions and rationale
evaluationandpromotion_evidence_link	provenance and evidence relationships
evaluationandpromotion_event_outbox	transactional Event publication
evaluationandpromotion_idempotency	duplicate-write protection
evaluationandpromotion_exception	exception, claim, dispute, or hold
evaluationandpromotion_correction	correction and supersession lineage
evaluationandpromotion_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

Online Monitoring and Drift

Metadata	Value
Volume	9
Chapter ID	V09-26
Subject	OnlineMonitoringandDrift
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Online Monitoring and Drift.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
OnlineMonitoringandDrift	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OnlineMonitoringandDriftVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OnlineMonitoringandDriftDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OnlineMonitoringandDriftEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
OnlineMonitoringandDriftException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
OnlineMonitoringandDrift {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateOnlineMonitoringandDrift
- ValidateOnlineMonitoringandDrift
- ApproveOnlineMonitoringandDrift
- ActivateOnlineMonitoringandDrift
- SuspendOnlineMonitoringandDrift
- CorrectOnlineMonitoringandDrift
- RetireOnlineMonitoringandDrift

Reference Events:

- OnlineMonitoringandDriftCreated
- OnlineMonitoringandDriftValidated
- OnlineMonitoringandDriftApproved
- OnlineMonitoringandDriftActivated
- OnlineMonitoringandDriftSuspended
- OnlineMonitoringandDriftCorrected
- OnlineMonitoringandDriftRetired

Reference queries:

- GetOnlineMonitoringandDrift
- ListOnlineMonitoringandDriftRecords
- GetOnlineMonitoringandDriftTimeline
- GetOnlineMonitoringandDriftEvidence
- GetOnlineMonitoringandDriftDecisionPackage
- GetOnlineMonitoringandDriftExceptions
- GetOnlineMonitoringandDriftAuditTrail

Reference API surface:

```

POST /api/v1/onlineMonitoringandDrift/commands
GET  /api/v1/onlineMonitoringandDrift/{id}
GET  /api/v1/onlineMonitoringandDrift/{id}/timeline
GET  /api/v1/onlineMonitoringandDrift/{id}/evidence
GET  /api/v1/onlineMonitoringandDrift/{id}/decisions
POST /api/v1/onlineMonitoringandDrift/{id}/exceptions
POST /api/v1/onlineMonitoringandDrift/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
onlinemonitoringanddrift_aggregate	authoritative subject
onlinemonitoringanddrift_version	immutable submitted, approved, issued, or signed versions
onlinemonitoringanddrift_relationship	governed relationships
onlinemonitoringanddrift_decision	binding Decisions and rationale
onlinemonitoringanddrift_evidence_link	provenance and evidence relationships
onlinemonitoringanddrift_event_outbox	transactional Event publication
onlinemonitoringanddrift_idempotency	duplicate-write protection
onlinemonitoringanddrift_exception	exception, claim, dispute, or hold
onlinemonitoringanddrift_correction	correction and supersession lineage
onlinemonitoringanddrift_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Kill Switch and Incident Containment

Metadata	Value
Volume	9
Chapter ID	V09-27
Subject	KillSwitchandIncidentContainment
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Kill Switch and Incident Containment.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
KillSwitchandIncidentContainment	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KillSwitchandIncidentContainmentVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KillSwitchandIncidentContainmentDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KillSwitchandIncidentContainmentEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
KillSwitchandIncidentContainmentException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

KillSwitchandIncidentContainment {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateKillSwitchandIncidentContainment
- ValidateKillSwitchandIncidentContainment
- ApproveKillSwitchandIncidentContainment
- ActivateKillSwitchandIncidentContainment
- SuspendKillSwitchandIncidentContainment
- CorrectKillSwitchandIncidentContainment
- RetireKillSwitchandIncidentContainment

Reference Events:

- KillSwitchandIncidentContainmentCreated
- KillSwitchandIncidentContainmentValidated
- KillSwitchandIncidentContainmentApproved
- KillSwitchandIncidentContainmentActivated
- KillSwitchandIncidentContainmentSuspended
- KillSwitchandIncidentContainmentCorrected
- KillSwitchandIncidentContainmentRetired

Reference queries:

- GetKillSwitchandIncidentContainment
- ListKillSwitchandIncidentContainmentRecords
- GetKillSwitchandIncidentContainmentTimeline
- GetKillSwitchandIncidentContainmentEvidence
- GetKillSwitchandIncidentContainmentDecisionPackage
- GetKillSwitchandIncidentContainmentExceptions
- GetKillSwitchandIncidentContainmentAuditTrail

Reference API surface:

```

POST /api/v1/killswitchandincidentcontainment/commands
GET  /api/v1/killswitchandincidentcontainment/{id}
GET  /api/v1/killswitchandincidentcontainment/{id}/timeline
GET  /api/v1/killswitchandincidentcontainment/{id}/evidence
GET  /api/v1/killswitchandincidentcontainment/{id}/decisions
POST /api/v1/killswitchandincidentcontainment/{id}/exceptions
POST /api/v1/killswitchandincidentcontainment/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
killswitchandincidentcontainment_aggregate	authoritative subject
killswitchandincidentcontainment_version	immutable submitted, approved, issued, or signed versions
killswitchandincidentcontainment_relationship	governed relationships
killswitchandincidentcontainment_decision	binding Decisions and rationale
killswitchandincidentcontainment_evidence_link	provenance and evidence relationships
killswitchandincidentcontainment_event_outbox	transactional Event publication
killswitchandincidentcontainment_idempotency	duplicate-write protection
killswitchandincidentcontainment_exception	exception, claim, dispute, or hold
killswitchandincidentcontainment_correction	correction and supersession lineage
killswitchandincidentcontainment_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

AI Security and Prompt Injection

Metadata	Value
Volume	9
Chapter ID	V09-28
Subject	AI Security and Prompt Injection
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Security and Prompt Injection.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AI Security and Prompt Injection	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AI Security and Prompt Injection Version	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AI Security and Prompt Injection Decision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AI Security and Prompt Injection Evidence Set	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AI Security and Prompt Injection Exception	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AI Security and Prompt Injection {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAISecurityandPromptInjection
- ValidateAISecurityandPromptInjection
- ApproveAISecurityandPromptInjection
- ActivateAISecurityandPromptInjection
- SuspendAISecurityandPromptInjection
- CorrectAISecurityandPromptInjection
- RetireAISecurityandPromptInjection

Reference Events:

- AISecurityandPromptInjectionCreated
- AISecurityandPromptInjectionValidated
- AISecurityandPromptInjectionApproved
- AISecurityandPromptInjectionActivated
- AISecurityandPromptInjectionSuspended
- AISecurityandPromptInjectionCorrected
- AISecurityandPromptInjectionRetired

Reference queries:

- GetAISecurityandPromptInjection
- ListAISecurityandPromptInjectionRecords
- GetAISecurityandPromptInjectionTimeline
- GetAISecurityandPromptInjectionEvidence
- GetAISecurityandPromptInjectionDecisionPackage
- GetAISecurityandPromptInjectionExceptions
- GetAISecurityandPromptInjectionAuditTrail

Reference API surface:

```

POST /api/v1/aisecurityandpromptinjection/commands
GET  /api/v1/aisecurityandpromptinjection/{id}
GET  /api/v1/aisecurityandpromptinjection/{id}/timeline
GET  /api/v1/aisecurityandpromptinjection/{id}/evidence
GET  /api/v1/aisecurityandpromptinjection/{id}/decisions
POST /api/v1/aisecurityandpromptinjection/{id}/exceptions
POST /api/v1/aisecurityandpromptinjection/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aisecurityandpromptinjection_aggregate	authoritative subject
aisecurityandpromptinjection_version	immutable submitted, approved, issued, or signed versions
aisecurityandpromptinjection_relationship	governed relationships
aisecurityandpromptinjection_decision	binding Decisions and rationale
aisecurityandpromptinjection_evidence_link	provenance and evidence relationships
aisecurityandpromptinjection_event_outbox	transactional Event publication
aisecurityandpromptinjection_idempotency	duplicate-write protection
aisecurityandpromptinjection_exception	exception, claim, dispute, or hold
aisecurityandpromptinjection_correction	correction and supersession lineage
aisecurityandpromptinjection_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Browser Agent Sandboxing

Metadata	Value
Volume	9
Chapter ID	V09-29
Subject	BrowserAgentSandboxing
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Browser Agent Sandboxing.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
BrowserAgentSandboxing	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
BrowserAgentSandboxingVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
BrowserAgentSandboxingDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
BrowserAgentSandboxingEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
BrowserAgentSandboxingException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
BrowserAgentSandboxing {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```



```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateBrowserAgentSandboxing
- ValidateBrowserAgentSandboxing
- ApproveBrowserAgentSandboxing
- ActivateBrowserAgentSandboxing
- SuspendBrowserAgentSandboxing
- CorrectBrowserAgentSandboxing
- RetireBrowserAgentSandboxing

Reference Events:

- BrowserAgentSandboxingCreated
- BrowserAgentSandboxingValidated
- BrowserAgentSandboxingApproved
- BrowserAgentSandboxingActivated
- BrowserAgentSandboxingSuspended
- BrowserAgentSandboxingCorrected
- BrowserAgentSandboxingRetired

Reference queries:

- GetBrowserAgentSandboxing
- ListBrowserAgentSandboxingRecords
- GetBrowserAgentSandboxingTimeline
- GetBrowserAgentSandboxingEvidence
- GetBrowserAgentSandboxingDecisionPackage
- GetBrowserAgentSandboxingExceptions
- GetBrowserAgentSandboxingAuditTrail

Reference API surface:

```

POST /api/v1/browseragentsandboxing/commands
GET  /api/v1/browseragentsandboxing/{id}
GET  /api/v1/browseragentsandboxing/{id}/timeline
GET  /api/v1/browseragentsandboxing/{id}/evidence
GET  /api/v1/browseragentsandboxing/{id}/decisions
POST /api/v1/browseragentsandboxing/{id}/exceptions
POST /api/v1/browseragentsandboxing/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
browseragentsandboxing_aggregate	authoritative subject
browseragentsandboxing_version	immutable submitted, approved, issued, or signed versions
browseragentsandboxing_relationship	governed relationships
browseragentsandboxing_decision	binding Decisions and rationale
browseragentsandboxing_evidence_link	provenance and evidence relationships
browseragentsandboxing_event_outbox	transactional Event publication
browseragentsandboxing_idempotency	duplicate-write protection
browseragentsandboxing_exception	exception, claim, dispute, or hold
browseragentsandboxing_correction	correction and supersession lineage
browseragentsandboxing_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Code and Data Tool Sandboxing

Metadata	Value
Volume	9
Chapter ID	V09-30
Subject	CodeandDataToolSandboxing
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Code and Data Tool Sandboxing**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
CodeandDataToolSandboxing	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CodeandDataToolSandboxingVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CodeandDataToolSandboxingDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CodeandDataToolSandboxingEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
CodeandDataToolSandboxingException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
CodeandDataToolSandboxing {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateCodeandDataToolSandboxing
- ValidateCodeandDataToolSandboxing
- ApproveCodeandDataToolSandboxing
- ActivateCodeandDataToolSandboxing
- SuspendCodeandDataToolSandboxing
- CorrectCodeandDataToolSandboxing
- RetireCodeandDataToolSandboxing

Reference Events:

- CodeandDataToolSandboxingCreated
- CodeandDataToolSandboxingValidated
- CodeandDataToolSandboxingApproved
- CodeandDataToolSandboxingActivated
- CodeandDataToolSandboxingSuspended
- CodeandDataToolSandboxingCorrected
- CodeandDataToolSandboxingRetired

Reference queries:

- GetCodeandDataToolSandboxing
- ListCodeandDataToolSandboxingRecords
- GetCodeandDataToolSandboxingTimeline
- GetCodeandDataToolSandboxingEvidence
- GetCodeandDataToolSandboxingDecisionPackage
- GetCodeandDataToolSandboxingExceptions
- GetCodeandDataToolSandboxingAuditTrail

Reference API surface:

```

POST /api/v1/codeanddatatoolsandboxing/commands
GET  /api/v1/codeanddatatoolsandboxing/{id}
GET  /api/v1/codeanddatatoolsandboxing/{id}/timeline
GET  /api/v1/codeanddatatoolsandboxing/{id}/evidence
GET  /api/v1/codeanddatatoolsandboxing/{id}/decisions
POST /api/v1/codeanddatatoolsandboxing/{id}/exceptions
POST /api/v1/codeanddatatoolsandboxing/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```



```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
codeanddatatoolsandboxing_aggregate	authoritative subject
codeanddatatoolsandboxing_version	immutable submitted, approved, issued, or signed versions
codeanddatatoolsandboxing_relationship	governed relationships
codeanddatatoolsandboxing_decision	binding Decisions and rationale
codeanddatatoolsandboxing_evidence_link	provenance and evidence relationships
codeanddatatoolsandboxing_event_outbox	transactional Event publication
codeanddatatoolsandboxing_idempotency	duplicate-write protection
codeanddatatoolsandboxing_exception	exception, claim, dispute, or hold
codeanddatatoolsandboxing_correction	correction and supersession lineage
codeanddatatoolsandboxing_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Cost and Budget Governance

Metadata	Value
Volume	9
Chapter ID	V09-31
Subject	AICostandBudgetGovernance
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Cost and Budget Governance**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AICostandBudgetGovernance	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AICostandBudgetGovernanceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AICostandBudgetGovernanceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AICostandBudgetGovernanceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AICostandBudgetGovernanceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AICostandBudgetGovernance {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAICostandBudgetGovernance
- ValidateAICostandBudgetGovernance
- ApproveAICostandBudgetGovernance
- ActivateAICostandBudgetGovernance
- SuspendAICostandBudgetGovernance
- CorrectAICostandBudgetGovernance
- RetireAICostandBudgetGovernance

Reference Events:

- AICostandBudgetGovernanceCreated
- AICostandBudgetGovernanceValidated
- AICostandBudgetGovernanceApproved
- AICostandBudgetGovernanceActivated
- AICostandBudgetGovernanceSuspended
- AICostandBudgetGovernanceCorrected
- AICostandBudgetGovernanceRetired

Reference queries:

- GetAICostandBudgetGovernance
- ListAICostandBudgetGovernanceRecords
- GetAICostandBudgetGovernanceTimeline
- GetAICostandBudgetGovernanceEvidence
- GetAICostandBudgetGovernanceDecisionPackage
- GetAICostandBudgetGovernanceExceptions
- GetAICostandBudgetGovernanceAuditTrail

Reference API surface:

```

POST /api/v1/aicostandbudgetgovernance/commands
GET  /api/v1/aicostandbudgetgovernance/{id}
GET  /api/v1/aicostandbudgetgovernance/{id}/timeline
GET  /api/v1/aicostandbudgetgovernance/{id}/evidence
GET  /api/v1/aicostandbudgetgovernance/{id}/decisions
POST /api/v1/aicostandbudgetgovernance/{id}/exceptions
POST /api/v1/aicostandbudgetgovernance/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aicostandbudgetgovernance_aggregate	authoritative subject
aicostandbudgetgovernance_version	immutable submitted, approved, issued, or signed versions
aicostandbudgetgovernance_relationship	governed relationships
aicostandbudgetgovernance_decision	binding Decisions and rationale
aicostandbudgetgovernance_evidence_link	provenance and evidence relationships
aicostandbudgetgovernance_event_outbox	transactional Event publication
aicostandbudgetgovernance_idempotency	duplicate-write protection
aicostandbudgetgovernance_exception	exception, claim, dispute, or hold
aicostandbudgetgovernance_correction	correction and supersession lineage
aicostandbudgetgovernance_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Human Handoff

Metadata	Value
Volume	9
Chapter ID	V09-32
Subject	HumanHandoff
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Human Handoff**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
HumanHandoff	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
HumanHandoffVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
HumanHandoffDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
HumanHandoffEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
HumanHandoffException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
HumanHandoff {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateHumanHandoff
- ValidateHumanHandoff
- ApproveHumanHandoff
- ActivateHumanHandoff
- SuspendHumanHandoff
- CorrectHumanHandoff
- RetireHumanHandoff

Reference Events:

- HumanHandoffCreated
- HumanHandoffValidated
- HumanHandoffApproved
- HumanHandoffActivated
- HumanHandoffSuspended
- HumanHandoffCorrected
- HumanHandoffRetired

Reference queries:

- GetHumanHandoff
- ListHumanHandoffRecords
- GetHumanHandoffTimeline
- GetHumanHandoffEvidence
- GetHumanHandoffDecisionPackage
- GetHumanHandoffExceptions
- GetHumanHandoffAuditTrail

Reference API surface:

```

POST /api/v1/humanhandoff/commands
GET  /api/v1/humanhandoff/{id}
GET  /api/v1/humanhandoff/{id}/timeline
GET  /api/v1/humanhandoff/{id}/evidence
GET  /api/v1/humanhandoff/{id}/decisions
POST /api/v1/humanhandoff/{id}/exceptions
POST /api/v1/humanhandoff/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
humanhandoff_aggregate	authoritative subject
humanhandoff_version	immutable submitted, approved, issued, or signed versions
humanhandoff_relationship	governed relationships
humanhandoff_decision	binding Decisions and rationale
humanhandoff_evidence_link	provenance and evidence relationships
humanhandoff_event_outbox	transactional Event publication
humanhandoff_idempotency	duplicate-write protection
humanhandoff_exception	exception, claim, dispute, or hold
humanhandoff_correction	correction and supersession lineage
humanhandoff_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Compliance Service

Metadata	Value
Volume	9
Chapter ID	V09-33
Subject	AIComplianceService
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Compliance Service.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIComplianceService	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIComplianceServiceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIComplianceServiceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIComplianceServiceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIComplianceServiceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIComplianceService {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIComplianceService
- ValidateAIComplianceService
- ApproveAIComplianceService
- ActivateAIComplianceService
- SuspendAIComplianceService
- CorrectAIComplianceService
- RetireAIComplianceService

Reference Events:

- AIComplianceServiceCreated
- AIComplianceServiceValidated
- AIComplianceServiceApproved
- AIComplianceServiceActivated
- AIComplianceServiceSuspended
- AIComplianceServiceCorrected
- AIComplianceServiceRetired

Reference queries:

- GetAIComplianceService
- ListAIComplianceServiceRecords
- GetAIComplianceServiceTimeline
- GetAIComplianceServiceEvidence
- GetAIComplianceServiceDecisionPackage
- GetAIComplianceServiceExceptions
- GetAIComplianceServiceAuditTrail

Reference API surface:

```

POST /api/v1/aicomplianceservice/commands
GET  /api/v1/aicomplianceservice/{id}
GET  /api/v1/aicomplianceservice/{id}/timeline
GET  /api/v1/aicomplianceservice/{id}/evidence
GET  /api/v1/aicomplianceservice/{id}/decisions
POST /api/v1/aicomplianceservice/{id}/exceptions
POST /api/v1/aicomplianceservice/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aicomplianceservice_aggregate	authoritative subject
aicomplianceservice_version	immutable submitted, approved, issued, or signed versions
aicomplianceservice_relationship	governed relationships
aicomplianceservice_decision	binding Decisions and rationale
aicomplianceservice_evidence_link	provenance and evidence relationships
aicomplianceservice_event_outbox	transactional Event publication
aicomplianceservice_idempotency	duplicate-write protection
aicomplianceservice_exception	exception, claim, dispute, or hold
aicomplianceservice_correction	correction and supersession lineage
aicomplianceservice_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Vision Service

Metadata	Value
Volume	9
Chapter ID	V09-34
Subject	AI Vision Service
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Vision Service.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIVisionService	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVisionServiceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVisionServiceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVisionServiceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVisionServiceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

AIVisionService {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIVisionService
- ValidateAIVisionService
- ApproveAIVisionService
- ActivateAIVisionService
- SuspendAIVisionService
- CorrectAIVisionService
- RetireAIVisionService

Reference Events:

- AIVisionServiceCreated
- AIVisionServiceValidated
- AIVisionServiceApproved
- AIVisionServiceActivated
- AIVisionServiceSuspended
- AIVisionServiceCorrected
- AIVisionServiceRetired

Reference queries:

- GetAIVisionService
- ListAIVisionServiceRecords
- GetAIVisionServiceTimeline
- GetAIVisionServiceEvidence
- GetAIVisionServiceDecisionPackage
- GetAIVisionServiceExceptions
- GetAIVisionServiceAuditTrail

Reference API surface:

```

POST /api/v1/aivision-service/commands
GET  /api/v1/aivision-service/{id}
GET  /api/v1/aivision-service/{id}/timeline
GET  /api/v1/aivision-service/{id}/evidence
GET  /api/v1/aivision-service/{id}/decisions
POST /api/v1/aivision-service/{id}/exceptions
POST /api/v1/aivision-service/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aivisionservice_aggregate	authoritative subject
aivisionservice_version	immutable submitted, approved, issued, or signed versions
aivisionservice_relationship	governed relationships
aivisionservice_decision	binding Decisions and rationale
aivisionservice_evidence_link	provenance and evidence relationships
aivisionservice_event_outbox	transactional Event publication
aivisionservice_idempotency	duplicate-write protection
aivisionservice_exception	exception, claim, dispute, or hold
aivisionservice_correction	correction and supersession lineage
aivisionservice_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

AI Voice Service

Metadata	Value
Volume	9
Chapter ID	V09-35
Subject	AIVoiceService
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Voice Service.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIVoiceService	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVoiceServiceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVoiceServiceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVoiceServiceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIVoiceServiceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```

AIVoiceService {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}

```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIVoiceService
- ValidateAIVoiceService
- ApproveAIVoiceService
- ActivateAIVoiceService
- SuspendAIVoiceService
- CorrectAIVoiceService
- RetireAIVoiceService

Reference Events:

- AIVoiceServiceCreated
- AIVoiceServiceValidated
- AIVoiceServiceApproved
- AIVoiceServiceActivated
- AIVoiceServiceSuspended
- AIVoiceServiceCorrected
- AIVoiceServiceRetired

Reference queries:

- GetAIVoiceService
- ListAIVoiceServiceRecords
- GetAIVoiceServiceTimeline
- GetAIVoiceServiceEvidence
- GetAIVoiceServiceDecisionPackage
- GetAIVoiceServiceExceptions
- GetAIVoiceServiceAuditTrail

Reference API surface:

```

POST /api/v1/aivoiceservice/commands
GET  /api/v1/aivoiceservice/{id}
GET  /api/v1/aivoiceservice/{id}/timeline
GET  /api/v1/aivoiceservice/{id}/evidence
GET  /api/v1/aivoiceservice/{id}/decisions
POST /api/v1/aivoiceservice/{id}/exceptions
POST /api/v1/aivoiceservice/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aivoiceservice_aggregate	authoritative subject
aivoiceservice_version	immutable submitted, approved, issued, or signed versions
aivoiceservice_relationship	governed relationships
aivoiceservice_decision	binding Decisions and rationale
aivoiceservice_evidence_link	provenance and evidence relationships
aivoiceservice_event_outbox	transactional Event publication
aivoiceservice_idempotency	duplicate-write protection
aivoiceservice_exception	exception, claim, dispute, or hold
aivoiceservice_correction	correction and supersession lineage
aivoiceservice_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI NLP Service

Metadata	Value
Volume	9
Chapter ID	V09-36
Subject	AINLPService
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI NLP Service.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AINLPService	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AINLPServiceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AINLPServiceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AINLPServiceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AINLPServiceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AINLPService {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```

```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAINLPService
- ValidateAINLPService
- ApproveAINLPService
- ActivateAINLPService
- SuspendAINLPService
- CorrectAINLPService
- RetireAINLPService

Reference Events:

- AINLPServiceCreated
- AINLPServiceValidated
- AINLPServiceApproved
- AINLPServiceActivated
- AINLPServiceSuspended
- AINLPServiceCorrected
- AINLPServiceRetired

Reference queries:

- GetAINLPService
- ListAINLPServiceRecords
- GetAINLPServiceTimeline
- GetAINLPServiceEvidence
- GetAINLPServiceDecisionPackage
- GetAINLPServiceExceptions
- GetAINLPServiceAuditTrail

Reference API surface:

```

POST /api/v1/ainlp/service/commands
GET  /api/v1/ainlp/service/{id}
GET  /api/v1/ainlp/service/{id}/timeline
GET  /api/v1/ainlp/service/{id}/evidence
GET  /api/v1/ainlp/service/{id}/decisions
POST /api/v1/ainlp/service/{id}/exceptions
POST /api/v1/ainlp/service/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
ainlp_service_aggregate	authoritative subject
ainlp_service_version	immutable submitted, approved, issued, or signed versions
ainlp_service_relationship	governed relationships
ainlp_service_decision	binding Decisions and rationale
ainlp_service_evidence_link	provenance and evidence relationships
ainlp_service_event_outbox	transactional Event publication
ainlp_service_idempotency	duplicate-write protection
ainlp_service_exception	exception, claim, dispute, or hold
ainlp_service_correction	correction and supersession lineage
ainlp_service_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Automation Service

Metadata	Value
Volume	9
Chapter ID	V09-37
Subject	AIAutomationService
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Automation Service.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;
- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;

- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIAutomationService	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIAutomationServiceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIAutomationServiceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIAutomationServiceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIAutomationServiceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIAutomationService {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
- require exact expected version where concurrent change matters;
- record actor, represented principal, mandate, Policy, reason, and Evidence;
- use an idempotency key for retried writes;
- publish a completed-fact Event after durable state change;
- expose Outcome Unknown for unresolved external completion;
- preserve original and corrected states in historical reconstruction.

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
    sourceVersionRefs[]
    knowledgeCutoff
    evidenceSetRef
}
```



```

policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIAutomationService
- ValidateAIAutomationService
- ApproveAIAutomationService
- ActivateAIAutomationService
- SuspendAIAutomationService
- CorrectAIAutomationService
- RetireAIAutomationService

Reference Events:

- AIAutomationServiceCreated
- AIAutomationServiceValidated
- AIAutomationServiceApproved
- AIAutomationServiceActivated
- AIAutomationServiceSuspended
- AIAutomationServiceCorrected
- AIAutomationServiceRetired

Reference queries:

- GetAIAutomationService
- ListAIAutomationServiceRecords
- GetAIAutomationServiceTimeline
- GetAIAutomationServiceEvidence
- GetAIAutomationServiceDecisionPackage
- GetAIAutomationServiceExceptions
- GetAIAutomationServiceAuditTrail

Reference API surface:

```

POST /api/v1/aiautomationservice/commands
GET  /api/v1/aiautomationservice/{id}
GET  /api/v1/aiautomationservice/{id}/timeline
GET  /api/v1/aiautomationservice/{id}/evidence
GET  /api/v1/aiautomationservice/{id}/decisions
POST /api/v1/aiautomationservice/{id}/exceptions
POST /api/v1/aiautomationservice/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]
F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]

```

```

G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aiautomation_service_aggregate	authoritative subject
aiautomation_service_version	immutable submitted, approved, issued, or signed versions
aiautomation_service_relationship	governed relationships
aiautomation_service_decision	binding Decisions and rationale
aiautomation_service_evidence_link	provenance and evidence relationships
aiautomation_service_event_outbox	transactional Event publication
aiautomation_service_idempotency	duplicate-write protection
aiautomation_service_exception	exception, claim, dispute, or hold
aiautomation_service_correction	correction and supersession lineage
aiautomation_service_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
stale-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW

Area	Classification
behavior proven in another repository-grounded chapter contradictory ownership or path claims	retain that chapter's evidence and classification CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED IMPLEMENTED_WITH_GAP	Executable source and relevant behavior were inspected. Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION REFERENCE_ARCHITECTURE PLANNED_NOT_IMPLEMENTED DOCUMENTATION_ONLY_CLAIM	Only part of the target capability is evidenced. Normative target behavior; no claim that it exists today. Explicitly planned but not evidenced as executable. Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION REQUIRES_REPOSITORY_REVIEW	Sources disagree and no final owner has been established. Necessary executable evidence has not yet been inspected.

AI Edge and Federated Services

Metadata	Value
Volume	9
Chapter ID	V09-38
Subject	AIEdgeandFederatedServices
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Edge and Federated Services.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIEdgeandFederatedServices	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIEdgeandFederatedServicesVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIEdgeandFederatedServicesDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIEdgeandFederatedServicesEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIEdgeandFederatedServicesException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIEdgeandFederatedServices {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIEdgeandFederatedServices
- ValidateAIEdgeandFederatedServices
- ApproveAIEdgeandFederatedServices
- ActivateAIEdgeandFederatedServices
- SuspendAIEdgeandFederatedServices
- CorrectAIEdgeandFederatedServices
- RetireAIEdgeandFederatedServices

Reference Events:

- AIEdgeandFederatedServicesCreated
- AIEdgeandFederatedServicesValidated
- AIEdgeandFederatedServicesApproved
- AIEdgeandFederatedServicesActivated
- AIEdgeandFederatedServicesSuspended
- AIEdgeandFederatedServicesCorrected
- AIEdgeandFederatedServicesRetired

Reference queries:

- GetAIEdgeandFederatedServices
- ListAIEdgeandFederatedServicesRecords
- GetAIEdgeandFederatedServicesTimeline
- GetAIEdgeandFederatedServicesEvidence
- GetAIEdgeandFederatedServicesDecisionPackage
- GetAIEdgeandFederatedServicesExceptions
- GetAIEdgeandFederatedServicesAuditTrail

Reference API surface:

```

POST /api/v1/aiedgeandfederatedservices/commands
GET  /api/v1/aiedgeandfederatedservices/{id}
GET  /api/v1/aiedgeandfederatedservices/{id}/timeline
GET  /api/v1/aiedgeandfederatedservices/{id}/evidence
GET  /api/v1/aiedgeandfederatedservices/{id}/decisions
POST /api/v1/aiedgeandfederatedservices/{id}/exceptions
POST /api/v1/aiedgeandfederatedservices/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```



```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aiedgeandfederatedservices_aggregate	authoritative subject
aiedgeandfederatedservices_version	immutable submitted, approved, issued, or signed versions
aiedgeandfederatedservices_relationship	governed relationships
aiedgeandfederatedservices_decision	binding Decisions and rationale
aiedgeandfederatedservices_evidence_link	provenance and evidence relationships
aiedgeandfederatedservices_event_outbox	transactional Event publication
aiedgeandfederatedservices_idempotency	duplicate-write protection
aiedgeandfederatedservices_exception	exception, claim, dispute, or hold
aiedgeandfederatedservices_correction	correction and supersession lineage
aiedgeandfederatedservices_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Marketplace

Metadata	Value
Volume	9
Chapter ID	V09-39
Subject	AIMarketplace
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Marketplace.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIMarketplace	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMarketplaceVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMarketplaceDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMarketplaceEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIMarketplaceException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIMarketplace {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIMarketplace
- ValidateAIMarketplace
- ApproveAIMarketplace
- ActivateAIMarketplace
- SuspendAIMarketplace
- CorrectAIMarketplace
- RetireAIMarketplace

Reference Events:

- AIMarketplaceCreated
- AIMarketplaceValidated
- AIMarketplaceApproved
- AIMarketplaceActivated
- AIMarketplaceSuspended
- AIMarketplaceCorrected
- AIMarketplaceRetired

Reference queries:

- GetAIMarketplace
- ListAIMarketplaceRecords
- GetAIMarketplaceTimeline
- GetAIMarketplaceEvidence
- GetAIMarketplaceDecisionPackage
- GetAIMarketplaceExceptions
- GetAIMarketplaceAuditTrail

Reference API surface:

```

POST /api/v1/aimarketplace/commands
GET  /api/v1/aimarketplace/{id}
GET  /api/v1/aimarketplace/{id}/timeline
GET  /api/v1/aimarketplace/{id}/evidence
GET  /api/v1/aimarketplace/{id}/decisions
POST /api/v1/aimarketplace/{id}/exceptions
POST /api/v1/aimarketplace/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aimarketplace_aggregate	authoritative subject
aimarketplace_version	immutable submitted, approved, issued, or signed versions
aimarketplace_relationship	governed relationships
aimarketplace_decision	binding Decisions and rationale
aimarketplace_evidence_link	provenance and evidence relationships
aimarketplace_event_outbox	transactional Event publication
aimarketplace_idempotency	duplicate-write protection
aimarketplace_exception	exception, claim, dispute, or hold
aimarketplace_correction	correction and supersession lineage
aimarketplace_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Observability and AIOps

Metadata	Value
Volume	9
Chapter ID	V09-40
Subject	AIobservabilityandAIOps
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for AI Observability and AIOps.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIObservabilityandAIOps	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIObservabilityandAIOpsVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIObservabilityandAIOpsDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIObservabilityandAIOpsEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIObservabilityandAIOpsException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIObservabilityandAIOps {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIIObservabilityandAIOps
- ValidateAIIObservabilityandAIOps
- ApproveAIIObservabilityandAIOps
- ActivateAIIObservabilityandAIOps
- SuspendAIIObservabilityandAIOps
- CorrectAIIObservabilityandAIOps
- RetireAIIObservabilityandAIOps

Reference Events:

- AIIObservabilityandAIOpsCreated
- AIIObservabilityandAIOpsValidated
- AIIObservabilityandAIOpsApproved
- AIIObservabilityandAIOpsActivated
- AIIObservabilityandAIOpsSuspended
- AIIObservabilityandAIOpsCorrected
- AIIObservabilityandAIOpsRetired

Reference queries:

- GetAIIObservabilityandAIOps
- ListAIIObservabilityandAIOpsRecords
- GetAIIObservabilityandAIOpsTimeline
- GetAIIObservabilityandAIOpsEvidence
- GetAIIObservabilityandAIOpsDecisionPackage
- GetAIIObservabilityandAIOpsExceptions
- GetAIIObservabilityandAIOpsAuditTrail

Reference API surface:

```

POST /api/v1/aiobservabilityandaiops/commands
GET  /api/v1/aiobservabilityandaiops/{id}
GET  /api/v1/aiobservabilityandaiops/{id}/timeline
GET  /api/v1/aiobservabilityandaiops/{id}/evidence
GET  /api/v1/aiobservabilityandaiops/{id}/decisions
POST /api/v1/aiobservabilityandaiops/{id}/exceptions
POST /api/v1/aiobservabilityandaiops/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aiobservabilityandaioops_aggregate	authoritative subject
aiobservabilityandaioops_version	immutable submitted, approved, issued, or signed versions
aiobservabilityandaioops_relationship	governed relationships
aiobservabilityandaioops_decision	binding Decisions and rationale
aiobservabilityandaioops_evidence_link	provenance and evidence relationships
aiobservabilityandaioops_event_outbox	transactional Event publication
aiobservabilityandaioops_idempotency	duplicate-write protection
aiobservabilityandaioops_exception	exception, claim, dispute, or hold
aiobservabilityandaioops_correction	correction and supersession lineage
aiobservabilityandaioops_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Predictive Analytics and Forecast ML

Metadata	Value
Volume	9
Chapter ID	V09-41
Subject	PredictiveAnalyticsandForecastML
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Predictive Analytics and Forecast ML.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
PredictiveAnalyticsandForecastML	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PredictiveAnalyticsandForecastMLVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PredictiveAnalyticsandForecastMLDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PredictiveAnalyticsandForecastMLEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
PredictiveAnalyticsandForecastMLException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
PredictiveAnalyticsandForecastML {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreatePredictiveAnalyticsandForecastML
- ValidatePredictiveAnalyticsandForecastML
- ApprovePredictiveAnalyticsandForecastML
- ActivatePredictiveAnalyticsandForecastML
- SuspendPredictiveAnalyticsandForecastML
- CorrectPredictiveAnalyticsandForecastML
- RetirePredictiveAnalyticsandForecastML

Reference Events:

- PredictiveAnalyticsandForecastMLCreated
- PredictiveAnalyticsandForecastMLValidated
- PredictiveAnalyticsandForecastMLApproved
- PredictiveAnalyticsandForecastMLActivated
- PredictiveAnalyticsandForecastMLSuspended
- PredictiveAnalyticsandForecastMLCorrected
- PredictiveAnalyticsandForecastMLRetired

Reference queries:

- GetPredictiveAnalyticsandForecastML
- ListPredictiveAnalyticsandForecastMLRecords
- GetPredictiveAnalyticsandForecastMLTimeline
- GetPredictiveAnalyticsandForecastMLEvidence
- GetPredictiveAnalyticsandForecastMLDecisionPackage
- GetPredictiveAnalyticsandForecastMLExceptions
- GetPredictiveAnalyticsandForecastMLAuditTrail

Reference API surface:

```

POST /api/v1/predictiveanalyticsandforecastml/commands
GET  /api/v1/predictiveanalyticsandforecastml/{id}
GET  /api/v1/predictiveanalyticsandforecastml/{id}/timeline
GET  /api/v1/predictiveanalyticsandforecastml/{id}/evidence
GET  /api/v1/predictiveanalyticsandforecastml/{id}/decisions
POST /api/v1/predictiveanalyticsandforecastml/{id}/exceptions
POST /api/v1/predictiveanalyticsandforecastml/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
predictiveanalyticsandforecastml_aggregate	authoritative subject
predictiveanalyticsandforecastml_version	immutable submitted, approved, issued, or signed versions
predictiveanalyticsandforecastml_relationship	governed relationships
predictiveanalyticsandforecastml_decision	binding Decisions and rationale
predictiveanalyticsandforecastml_evidence_link	provenance and evidence relationships
predictiveanalyticsandforecastml_event_outbox	transactional Event publication
predictiveanalyticsandforecastml_idempotency	duplicate-write protection
predictiveanalyticsandforecastml_exception	exception, claim, dispute, or hold
predictiveanalyticsandforecastml_correction	correction and supersession lineage
predictiveanalyticsandforecastml_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Fraud Engine

Metadata	Value
Volume	9
Chapter ID	V09-42
Subject	FraudEngine
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Fraud Engine.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
FraudEngine	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
FraudEngineVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
FraudEngineDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
FraudEngineEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
FraudEngineException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
FraudEngine {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateFraudEngine
- ValidateFraudEngine
- ApproveFraudEngine
- ActivateFraudEngine
- SuspendFraudEngine
- CorrectFraudEngine
- RetireFraudEngine

Reference Events:

- FraudEngineCreated
- FraudEngineValidated
- FraudEngineApproved
- FraudEngineActivated
- FraudEngineSuspended
- FraudEngineCorrected
- FraudEngineRetired

Reference queries:

- GetFraudEngine
- ListFraudEngineRecords
- GetFraudEngineTimeline
- GetFraudEngineEvidence
- GetFraudEngineDecisionPackage
- GetFraudEngineExceptions
- GetFraudEngineAuditTrail

Reference API surface:

```

POST /api/v1/fraudengine/commands
GET  /api/v1/fraudengine/{id}
GET  /api/v1/fraudengine/{id}/timeline
GET  /api/v1/fraudengine/{id}/evidence
GET  /api/v1/fraudengine/{id}/decisions
POST /api/v1/fraudengine/{id}/exceptions
POST /api/v1/fraudengine/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
fraudengine_aggregate	authoritative subject
fraudengine_version	immutable submitted, approved, issued, or signed versions
fraudengine_relationship	governed relationships
fraudengine_decision	binding Decisions and rationale
fraudengine_evidence_link	provenance and evidence relationships
fraudengine_event_outbox	transactional Event publication
fraudengine_idempotency	duplicate-write protection
fraudengine_exception	exception, claim, dispute, or hold
fraudengine_correction	correction and supersession lineage
fraudengine_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Search and Ranking

Metadata	Value
Volume	9
Chapter ID	V09-43
Subject	SearchandRanking
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Search and Ranking.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
SearchandRanking	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SearchandRankingVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SearchandRankingDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SearchandRankingEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
SearchandRankingException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
SearchandRanking {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateSearchandRanking
- ValidateSearchandRanking
- ApproveSearchandRanking
- ActivateSearchandRanking
- SuspendSearchandRanking
- CorrectSearchandRanking
- RetireSearchandRanking

Reference Events:

- SearchandRankingCreated
- SearchandRankingValidated
- SearchandRankingApproved
- SearchandRankingActivated
- SearchandRankingSuspended
- SearchandRankingCorrected
- SearchandRankingRetired

Reference queries:

- GetSearchandRanking
- ListSearchandRankingRecords
- GetSearchandRankingTimeline
- GetSearchandRankingEvidence
- GetSearchandRankingDecisionPackage
- GetSearchandRankingExceptions
- GetSearchandRankingAuditTrail

Reference API surface:

```

POST /api/v1/searchandraking/commands
GET  /api/v1/searchandraking/{id}
GET  /api/v1/searchandraking/{id}/timeline
GET  /api/v1/searchandraking/{id}/evidence
GET  /api/v1/searchandraking/{id}/decisions
POST /api/v1/searchandraking/{id}/exceptions
POST /api/v1/searchandraking/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
searchandraking_aggregate	authoritative subject
searchandraking_version	immutable submitted, approved, issued, or signed versions
searchandraking_relationship	governed relationships
searchandraking_decision	binding Decisions and rationale
searchandraking_evidence_link	provenance and evidence relationships
searchandraking_event_outbox	transactional Event publication
searchandraking_idempotency	duplicate-write protection
searchandraking_exception	exception, claim, dispute, or hold
searchandraking_correction	correction and supersession lineage
searchandraking_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

Domain AI Role Catalogue

Metadata	Value
Volume	9
Chapter ID	V09-44
Subject	DomainAIRoleCatalogue
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **define the ownership, lifecycle, policy, evidence, Tools, security, evaluation, failure, recovery, and repository verification requirements for Domain AI Role Catalogue.**

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
DomainAIRoleCatalogue	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DomainAIRoleCatalogueVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DomainAIRoleCatalogueDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DomainAIRoleCatalogueEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
DomainAIRoleCatalogueException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
DomainAIRoleCatalogue {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```

```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateDomainAIRoleCatalogue
- ValidateDomainAIRoleCatalogue
- ApproveDomainAIRoleCatalogue
- ActivateDomainAIRoleCatalogue
- SuspendDomainAIRoleCatalogue
- CorrectDomainAIRoleCatalogue
- RetireDomainAIRoleCatalogue

Reference Events:

- DomainAIRoleCatalogueCreated
- DomainAIRoleCatalogueValidated
- DomainAIRoleCatalogueApproved
- DomainAIRoleCatalogueActivated
- DomainAIRoleCatalogueSuspended
- DomainAIRoleCatalogueCorrected
- DomainAIRoleCatalogueRetired

Reference queries:

- GetDomainAIRoleCatalogue
- ListDomainAIRoleCatalogueRecords
- GetDomainAIRoleCatalogueTimeline
- GetDomainAIRoleCatalogueEvidence
- GetDomainAIRoleCatalogueDecisionPackage
- GetDomainAIRoleCatalogueExceptions
- GetDomainAIRoleCatalogueAuditTrail

Reference API surface:

```

POST /api/v1/domainairolecatalogue/commands
GET  /api/v1/domainairolecatalogue/{id}
GET  /api/v1/domainairolecatalogue/{id}/timeline
GET  /api/v1/domainairolecatalogue/{id}/evidence
GET  /api/v1/domainairolecatalogue/{id}/decisions
POST /api/v1/domainairolecatalogue/{id}/exceptions
POST /api/v1/domainairolecatalogue/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
domainairolecatalogue_aggregate	authoritative subject
domainairolecatalogue_version	immutable submitted, approved, issued, or signed versions
domainairolecatalogue_relationship	governed relationships
domainairolecatalogue_decision	binding Decisions and rationale
domainairolecatalogue_evidence_link	provenance and evidence relationships
domainairolecatalogue_event_outbox	transactional Event publication
domainairolecatalogue_idempotency	duplicate-write protection
domainairolecatalogue_exception	exception, claim, dispute, or hold
domainairolecatalogue_correction	correction and supersession lineage
domainairolecatalogue_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- plan and execute only through typed registered Tools
- retain source context and knowledge cutoff
- expose confidence, limitations and policy result
- observe outcomes and trigger recovery or human handoff

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- central provider and Tool governance
- prompt injection and data exfiltration defense
- sandbox and egress policy
- budget, approval, kill switch and incident containment

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.

AI Canonical Ownership Matrix

Metadata	Value
Volume	9
Chapter ID	V09-SYNTH
Subject	AIOwnershipMatrix
Status	Generated First-Draft Architecture Skeleton - Domain-Specific Expansion Required
Content maturity	TEMPLATE_DERIVED_REFERENCE_SKELETON unless repository citations prove otherwise
Default implementation classification	REFERENCE_ARCHITECTURE / REQUIRES_REPOSITORY_REVIEW
Accountable owner	AI Architecture Council
Evidence rule	No implementation claim without inspected executable evidence
Repository	chinair88-prog/allinb2c

1. Executive Orientation

This chapter defines how GTOS shall **assign AI Orchestrator, AI Runtime, MCP, Workflow, Event Bus, Domain services, memory, RAG, providers, evaluation, policy, Tools, Portals, and operators without shadow ownership**.

The specification separates reality, observations, knowledge, Decisions, Intent, execution, and outcomes. A component name, database row, UI label, AI answer, or workflow state does not acquire authority merely because it is visible or technically convenient.

Reality or external outcome

- ≠ Observation
- ≠ Claim
- ≠ derived Perspective
- ≠ Prediction
- ≠ Recommendation
- ≠ authorized Decision
- ≠ execution Intent

1.1 Accountable ownership

The accountable owner is **AI Architecture Council**. Supporting applications, shared infrastructure, integration adapters, analytics, and AI coordinate or present the capability but shall not silently own its authoritative propositions.

1.2 Scope

- create and identify the subject
- validate evidence and policy
- make authorized Decisions
- execute typed Commands
- publish completed-fact Events
- query current and historical perspectives
- correct and reconcile outcomes

1.3 Non-goals

- treating a registered Gradle module, catalog record, README, frontend route, migration plan, or generated report as proof of end-to-end implementation;
- direct cross-Domain database writes;

- generic status values that combine lifecycle, approval, provider, document, payment, dispute, and correction state;
- silent replacement of history;
- AI-created authority or fabricated external outcomes.

2. Ubiquitous Language and Subject Model

Subject	Required interpretation
AIOwnershipMatrix	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOwnershipMatrixVersion	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOwnershipMatrixDecision	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOwnershipMatrixEvidenceSet	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit
AIOwnershipMatrixException	Independent identified subject, governed subordinate entity, or qualified perspective; ownership must be explicit

Reference aggregate:

```
AIOwnershipMatrix {
  subjectId
  tenantId
  organizationContext
  sourceIdentityRefs[]
  sourceVersionRefs[]
  lifecycleState
  relationshipRefs[]
  decisionRefs[]
  intentRefs[]
  evidenceSetRef
  knowledgeCutoff
  policyRefs[]
  authorityRef
  jurisdictionRefs[]
  createdAt
  effectiveAt
  updatedAt
  version
  correctionOf?
  correlationId
}
```

2.1 Core invariants

- Stable identity precedes relationship, state, and cross-system correlation.
- Effective time and recorded time remain distinct where business truth changes over time.
- Recommendation, Decision, Intent, execution, and outcome are separate concepts.
- No Portal, report, search index, cache, workflow, or AI Agent becomes a shadow owner of Domain truth.
- Binding submitted, approved, issued, accepted, or signed versions are immutable.
- Corrections append governed facts and preserve the original record.
- External authority and provider outcomes retain their source and qualification.
- Tenant and represented-principal context come from trusted authentication and mandate evaluation.
- Retried writes are idempotent and concurrency guarded.
- Outcome Unknown is explicit and reconciled before risky duplicate execution.

2.2 Required standards

- stable identifiers
- typed contracts
- immutable binding versions
- transactional outbox
- idempotent Commands
- explicit Outcome Unknown
- correction lineage
- machine-verifiable conformance

3. Actors, Roles, and Authority

A conformant design distinguishes:

- the natural person acting;
- the service or AI identity invoking a Tool;
- the organization or principal represented;
- the role or relationship granting context;
- the mandate or delegation granting action authority;
- the Policy and jurisdiction limiting the action;
- the reviewer or approver responsible for a binding Decision.

Authority is evaluated at action time. Authentication alone does not grant business authority. Cached permissions and browser-supplied tenant headers are insufficient for high-impact actions.

4. Lifecycle and State Model

State	Semantics
PROPOSED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
VALIDATION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
APPROVAL_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
ACTIVE	Distinct governed condition with entry, exit, timeout, correction and authority semantics
SUSPENDED	Distinct governed condition with entry, exit, timeout, correction and authority semantics
CORRECTION_PENDING	Distinct governed condition with entry, exit, timeout, correction and authority semantics
RETIRED	Distinct governed condition with entry, exit, timeout, correction and authority semantics

```
stateDiagram-v2
    [*] --> PROPOSED
    PROPOSED --> VALIDATION_PENDING
    VALIDATION_PENDING --> APPROVAL_PENDING
    APPROVAL_PENDING --> ACTIVE
    ACTIVE --> SUSPENDED
    SUSPENDED --> CORRECTION_PENDING
    CORRECTION_PENDING --> RETIRED
```

The diagram is a primary reading path. Real implementations may add parallel review, amendment, appeal, timeout, suspension, partial completion, reconciliation, correction, and terminal failure without collapsing their meaning.

4.1 Transition rules

- transition through a typed, authorized Command;
 - require exact expected version where concurrent change matters;
 - record actor, represented principal, mandate, Policy, reason, and Evidence;
 - use an idempotency key for retried writes;
 - publish a completed-fact Event after durable state change;
 - expose Outcome Unknown for unresolved external completion;
 - preserve original and corrected states in historical reconstruction.
-

5. Decision, Evidence, and Knowledge Cutoff

Reference Decision:

```
Decision {
    decisionId
    decisionType
    subjectRef
}
```



```

sourceVersionRefs[]
knowledgeCutoff
evidenceSetRef
policyVersionRefs[]
recommendationRefs[]
principalId
representedOrganizationId
mandateRef
rationale
outcome
decidedAt
effectiveAt
correctionOf?
}

```

Evidence records identify issuer or observer, subject, version, observed time, received time, effective interval, integrity, confidence, qualification, contradiction, correction, retention, residency, and access class. A recommendation may assist the Decision but cannot promote itself to the Decision.

6. Commands, Events, Queries, and Contracts

Reference Commands:

- CreateAIOwnershipMatrix
- ValidateAIOwnershipMatrix
- ApproveAIOwnershipMatrix
- ActivateAIOwnershipMatrix
- SuspendAIOwnershipMatrix
- CorrectAIOwnershipMatrix
- RetireAIOwnershipMatrix

Reference Events:

- AIOwnershipMatrixCreated
- AIOwnershipMatrixValidated
- AIOwnershipMatrixApproved
- AIOwnershipMatrixActivated
- AIOwnershipMatrixSuspended
- AIOwnershipMatrixCorrected
- AIOwnershipMatrixRetired

Reference queries:

- GetAIOwnershipMatrix
- ListAIOwnershipMatrixRecords
- GetAIOwnershipMatrixTimeline
- GetAIOwnershipMatrixEvidence
- GetAIOwnershipMatrixDecisionPackage
- GetAIOwnershipMatrixExceptions
- GetAIOwnershipMatrixAuditTrail

Reference API surface:

```

POST /api/v1/aiownershipmatrix/commands
GET  /api/v1/aiownershipmatrix/{id}
GET  /api/v1/aiownershipmatrix/{id}/timeline
GET  /api/v1/aiownershipmatrix/{id}/evidence
GET  /api/v1/aiownershipmatrix/{id}/decisions
POST /api/v1/aiownershipmatrix/{id}/exceptions
POST /api/v1/aiownershipmatrix/{id}/corrections

```

These endpoints are reference contracts, not claims that exact routes exist.

7. Workflow, Failure, and Compensation

flowchart TD

```

A[Validate identity version and authority] --> B[Collect required evidence]
B --> C{Policy and evidence gates pass?}
C -- No --> D[Open exception or request correction]
D --> B
C -- Yes --> E[Record Decision or execution Intent]
E --> F[Execute Domain command or external request]

```

```

F --> G{Outcome known?}
G -- Success --> H[Record fact and publish Event]
G -- Failure --> I[Record failure and compensation options]
G -- Unknown --> J[Enter Outcome Unknown]
J --> K[Query source or wait for verified callback]
K --> G
H --> L[Update projections and downstream workflow]

```

Representative failure modes:

- stale or wrong source version
- duplicate submission or replay
- external timeout after possible success
- contradictory source observations
- lost Event after committed state
- approval or mandate revoked during execution
- downstream action before a mandatory gate
- partial completion across multiple subjects
- evidence later retracted or corrected
- tenant, facility, channel or jurisdiction mismatch

Recovery shall query the authoritative source before duplicate execution, preserve partial success, use compensating Intent rather than destructive rewrite, publish corrections, quarantine high-impact ambiguity, and record affected downstream subjects.

8. Data and Persistence Architecture

Reference table families:

Table family	Purpose
aiownershipmatrix_aggregate	authoritative subject
aiownershipmatrix_version	immutable submitted, approved, issued, or signed versions
aiownershipmatrix_relationship	governed relationships
aiownershipmatrix_decision	binding Decisions and rationale
aiownershipmatrix_evidence_link	provenance and evidence relationships
aiownershipmatrix_event_outbox	transactional Event publication
aiownershipmatrix_idempotency	duplicate-write protection
aiownershipmatrix_exception	exception, claim, dispute, or hold
aiownershipmatrix_correction	correction and supersession lineage
aiownershipmatrix_audit	privileged action evidence

The owning service controls schema and migration. Cross-Domain references use stable identifiers or contracts. Analytics, reports, caches, search, and vector indexes remain projections. Backup, restore, retention, legal hold, privacy, residency, and decommissioning requirements are explicit.

9. Portal and Experience Requirements

- Primary participant Portal: initiate, review exact versions, supply evidence, and respond to required actions.
- Operator Portal: inspect exceptions, correlation, source responses, retries, compensation, and downstream impact.
- Executive or oversight view: consume governed metrics without becoming a source of Domain truth.

Every Portal displays exact version, owner, state, freshness, required action, and uncertainty. It supports loading, empty, partial, stale, denied, error, timeout, and Outcome Unknown states; accessibility; locale, currency, unit, date, and RTL/LTR correctness; evidence audit; safe retry; and correction history.

10. AI Participation and Tool Governance

Permitted AI tasks:

- extract and classify evidence
- detect anomalies and missing fields
- summarize timelines and contradictions
- recommend next actions with sources and uncertainty

AI shall not own Domain records, grant itself authority, suppress contradictions, silently edit approved versions, fabricate external outcomes, or execute high-impact actions outside typed Tools, policy, approval, idempotency, sandbox, and outcome observation.

Every Agent Run records source context, prompt/model/tool versions, structured output, confidence and limitations, Policy result, approval, Tool calls, provider responses, final outcome, and recovery actions.

11. Security, Privacy, and Trust Boundaries

- least privilege and separation of duties
- tenant, organization, facility, channel and jurisdiction isolation
- integrity protection for binding evidence
- verified external callbacks and anti-replay controls

Additional controls include secret isolation, callback signature verification, audit of privileged action, emergency-access review, sensitive-data minimization, purpose limitation, and threat models for impersonation, tampering, replay, confused deputy, fraud, prompt injection, and exfiltration.

12. Reliability, SLOs, and Operations

SLI	Reference objective
authoritative read availability	99.9% monthly unless a stricter tier applies
acknowledged Command durability	no acknowledged write may be silently lost
Event publication	99.9% within five minutes after committed fact
duplicate prevention	100% for same tenant, Command class, and idempotency key
binding Decision audit completeness	100%
state-state detection	bounded by business deadline

Dashboards and runbooks cover backlog, error budget, dependency failure, retry budget, DLQ, workflow stall, provider outage, reconciliation, capacity, backup/restore, release/rollback, incident command, and post-incident correction.

13. Testing and Continuous Conformance

- aggregate invariant and transition tests
- authorization, delegation, separation-of-duties and tenant-isolation tests
- idempotency and optimistic-concurrency tests
- database migration, key, constraint and rollback tests
- contract, outbox/inbox, duplicate, replay and DLQ tests
- workflow restart, timeout and compensation tests
- Portal E2E tests for success, partial, stale, denied, error and Outcome Unknown states
- accessibility, locale, currency, unit and RTL/LTR tests
- security adversarial, fraud and data-exfiltration tests
- AI evaluation, prompt-injection, policy-bypass and unauthorized-tool tests
- historical reconstruction, correction and audit-completeness tests
- performance, capacity, recovery and operational exercise tests

Release evidence links each invariant and control to the test, environment, input data, result, owner, and artifact proving it. Generated test counts never replace semantic coverage.

14. Repository Review Boundary

- This chapter is a repository-review specification. Exact executable implementation has not been inspected for every target behavior.
- Registered modules, catalog entries, README claims, frontend contracts, or plans are not proof of complete implementation.

Area	Classification
target aggregate, API, tables, state and workflow in this chapter	REFERENCE_ARCHITECTURE
registered module or catalog claim	DOCUMENTATION_ONLY_CLAIM until source inspection
uninspected integration	REQUIRES_REPOSITORY_REVIEW
behavior proven in another repository-grounded chapter	retain that chapter's evidence and classification
contradictory ownership or path claims	CONFLICT_REQUIRES_CANONICALIZATION

No statement upgrades implementation status merely because a file, directory, module name, frontend contract, database entry, plan, or report exists.

15. Worked Conformance Scenario

An authorized principal initiates a Command against an exact subject version. The service validates identity, tenant, mandate, Policy, Evidence, and idempotency. It records the state transition and outbox entry atomically. If an external dependency participates, timeout produces Outcome Unknown; reconciliation queries the source or waits for a verified callback instead of repeating the action blindly.

Portals show source, freshness, version, required action, and uncertainty. AI may extract, compare, summarize, detect anomalies, and recommend through typed Tools, but cannot own the subject or fabricate success. Later correction appends new facts, identifies affected projections and workflows, and preserves the historical Decision context.

16. Conformance Checklist

- ☐ stable identity and proposition ownership are explicit;
- ☐ authority and mandate are evaluated at action time;
- ☐ state, Recommendation, Decision, Intent, execution, and outcome remain distinct;
- ☐ Commands are typed, authorized, idempotent, and version guarded;
- ☐ Events describe completed facts;
- ☐ Evidence includes provenance, time, integrity, qualification, and correction;
- ☐ external outcomes retain their authority source;
- ☐ Portals show freshness and uncertainty;
- ☐ AI is bounded by typed Tools, Policy, approval, and outcome observation;
- ☐ security, tenant, facility, channel, and jurisdiction boundaries are enforced;
- ☐ Outcome Unknown and reconciliation exist where needed;
- ☐ corrections preserve history;
- ☐ tests prove invariants and recovery;
- ☐ implementation status is not upgraded without executable evidence.

17. Status Vocabulary

Status	Meaning
VERIFIED_IMPLEMENTED	Executable source and relevant behavior were inspected.
IMPLEMENTED_WITH_GAP	Executable behavior exists with a material governance or completeness gap.
PARTIAL_IMPLEMENTATION	Only part of the target capability is evidenced.
REFERENCE_ARCHITECTURE	Normative target behavior; no claim that it exists today.
PLANNED_NOT_IMPLEMENTED	Explicitly planned but not evidenced as executable.
DOCUMENTATION_ONLY_CLAIM	Documentation or client contract claims behavior without verified backend evidence.
CONFLICT_REQUIRES_CANONICALIZATION	Sources disagree and no final owner has been established.
REQUIRES_REPOSITORY_REVIEW	Necessary executable evidence has not yet been inspected.